

Chapter 20.3 Dijkstra's Algorithm & Chapter 34 NP-Completeness

Joon Soo Yoo

June 10, 2025

Assignment

34

- ▶ Read §22.3, §34.1
- ▶ Problems
 - ▶ §22.3-2

Chapter 22: Single-Source Shortest Paths

BFS



- ▶ Chapter 22.1: The Bellman-Ford Algorithm
- ▶ Chapter 22.2: Single-Source Shortest Paths in Directed Acyclic Graphs
- ▶ **Chapter 22.3: Dijkstra's Algorithm**
- ▶ Chapter 22.4: Difference Constraints and Shortest Paths
- ▶ Chapter 22.5: Proofs of Shortest-Paths Properties

Relaxation in Shortest-Paths Algorithms

Core Concept: Relaxation


- ▶ Each vertex v maintains a **shortest-path estimate** $v.d$
 - ▶ $v.d$ is an upper bound on the weight of the shortest path from source s to v
- ▶ The goal is to iteratively **reduce** $v.d$ to the correct shortest-path weight $\delta(s, v)$

Initialization:

INITIALIZE-SINGLE-SOURCE(G, s)

```
1  for each vertex  $v \in G.V$ 
2       $v.d = \infty$ 
3       $v.\pi = \text{NIL}$ 
4   $s.d = 0$ 
```

The RELAX Operation



Relaxing an edge (u, v) with weight $w(u, v)$:

RELAX(u, v, w)

- 1 **if** $v.d > u.d + w(u, v)$
- 2 $v.d = u.d + w(u, v)$
- 3 $v.\pi = u$

Effect:

- ▶ Updates $v.d$ if a better path through u is found
- ▶ Updates $v.\pi$ to point to u


Dijkstra's Algorithm: Overview



Goal: Find the shortest paths from a source vertex s to all other vertices in a weighted, directed graph $G = (V, E)$, with $w(u, v) \geq 0$.

Key Concepts:

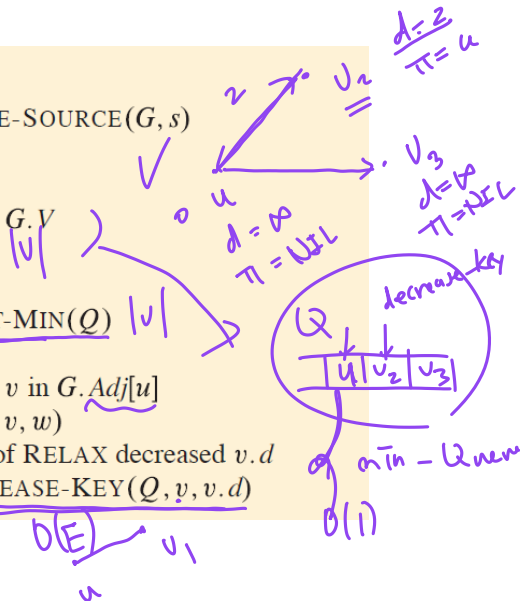
- ▶ Generalizes **BFS** to weighted graphs
- ▶ Uses a **greedy strategy** to iteratively build a shortest-paths tree
- ▶ Maintains:
 - ▶ Set S : vertices whose shortest-path weights are known
 - ▶ Min-priority queue $Q = V - S$









Dijkstra's Algorithm: Pseudocode

DIJKSTRA(G, w, s)

```
1  INITIALIZE-SINGLE-SOURCE( $G, s$ )
2   $S = \emptyset$ 
3   $Q = \emptyset$ 
4  for each vertex  $u \in G.V$ 
5    INSERT( $Q, u$ )
6  while  $Q \neq \emptyset$ 
7     $u = \text{EXTRACT-MIN}(Q)$ 
8     $S = S \cup \{u\}$ 
9    for each vertex  $v$  in  $G.Adj[u]$ 
10     RELAX( $u, v, w$ )
11     if the call of RELAX decreased  $v.d$ 
12       DECREASE-KEY( $Q, v, v.d$ )
```



Dijkstra's Algorithm: Diagram



Dijkstra's Algorithm: Step-by-Step (I)

Input: Directed graph $G = (V, E)$ with $w(u, v) \geq 0$ for all $(u, v) \in E$ and source s

Initialization:


1. For each $v \in V$: set $v.d \leftarrow \infty$, $v.\pi \leftarrow \text{NIL}$; then set $s.d \leftarrow 0$
2. Set $S \leftarrow \emptyset$ (vertices with finalized shortest-path weights)
3. Insert all $v \in V$ into a min-priority queue Q , keyed by $v.d$

Main Loop:

- ▶ While $Q \neq \emptyset$:
 - ▶ $u \leftarrow \text{EXTRACT-MIN}(Q)$
 - ▶ Add u to S
 - ▶ For each $v \in \text{Adj}[u]$, perform $\text{RELAX}(u, v, w)$
 - ▶ If $v.d$ changes, call $\text{DECREASE-KEY}(Q, v, v.d)$



Dijkstra's Algorithm: Key Properties (II)



Greedy Strategy:

- ▶ Always selects vertex $u \in V - S$ with the smallest $u.d$
- ▶ Once u is added to S , $u.d = \delta(s, u)$ is guaranteed

Invariant:

- ▶ $Q = V - S$ at the start of each iteration
- ▶ Each vertex is extracted from Q exactly once
- ▶ Each edge (u, v) is relaxed at most once

Output: For all $v \in V$, $v.d = \delta(s, v)$ and $v.\pi$ gives the predecessor on the shortest path

Correctness of Dijkstra's Algorithm

BFS

Theorem (Theorem 22.6)

Let $G = (V, E)$ be a weighted, directed graph with nonnegative edge weights $w(u, v) \geq 0$ and source vertex $s \in V$.

Then, after Dijkstra's algorithm terminates, for every vertex $u \in V$, we have:

$$u.d = \delta(s, u)$$

u
 $u.d = \delta(s, u)$

That is, the algorithm correctly computes the shortest path distance from s to u .

Corollary (Corollary 22.7)

The predecessor subgraph G_π forms a shortest-paths tree rooted at s .

Proof Sketch of Theorem 22.6 (1/2)

Strategy: Use **induction** on the number of iterations of the while loop (i.e., size of S).

Base case:

- ▶ When $S = \emptyset$, no vertex is added, so the claim trivially holds.
- ▶ When $S = \{s\}$, we have $s.d = 0 = \delta(s, s)$.

Inductive step:

- ▶ Assume for all $v \in S$, $v.d = \delta(s, v)$.
- ▶ Let u be the next vertex extracted from $Q = V - S$ (minimum d).
- ▶ Show: $u.d = \delta(s, u)$.

Key idea:

- ▶ Let y be the first vertex on a shortest path from s to u that is not in S .
- ▶ Let $x \in S$ be the predecessor of y on this path.

Proof Sketch of Theorem 22.6 (2/2)

Argument:

- ▶ From the inductive hypothesis: $x.d = \delta(s, x)$.
- ▶ Relaxation of edge (x, y) ensures: $y.d = \delta(s, y)$ at the time x was added to S .
- ▶ Because u was chosen as $\text{EXTRACT-MIN}(Q)$ and y is on the path to u , we have:

$$\delta(s, y) \leq \delta(s, u) \leq u.d \leq y.d$$

and since $y.d = \delta(s, y)$, all inequalities collapse:

$$u.d = \delta(s, u)$$

Conclusion: By induction, when the algorithm terminates, $u.d = \delta(s, u)$ for all $u \in V$.

Time Complexity of Dijkstra's Algorithm (Simple Array Implementation)

Assumption: Min-priority queue implemented using an unsorted array.

Operation Costs:

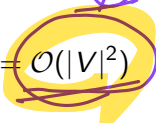
- ▶ $\text{INSERT}(v): \mathcal{O}(1)$
- ▶ $\text{EXTRACT-MIN}(): \mathcal{O}(|V|)$
- ▶ $\text{DECREASE-KEY}(v): \mathcal{O}(1)$



Total Number of Calls:

- ▶ $|V|$ calls to INSERT and EXTRACT-MIN
- ▶ At most $|E|$ calls to DECREASE-KEY (\because every edge relaxes once)

Total Time:

$$\mathcal{O}(|V| \cdot |V| + |E| \cdot 1) = \mathcal{O}(|V|^2 + |E|) = \mathcal{O}(|V|^2)$$


Use case: Acceptable for dense graphs

Time Complexity of Dijkstra's Algorithm (Binary Min-Heap Implementation)

min-heap

Assumption: Min-priority queue implemented using a *binary min-heap* with vertex-to-index mapping.

Operation Costs:

- ▶ INSERT(v): $\mathcal{O}(\log |V|)$ — insert into heap
- ▶ EXTRACT-MIN(): $\mathcal{O}(\log |V|)$ — remove min and heapify
- ▶ DECREASE-KEY(v): $\mathcal{O}(\log |V|)$ — heapify



$\mathcal{O}(\log n)$

Total Number of Calls:

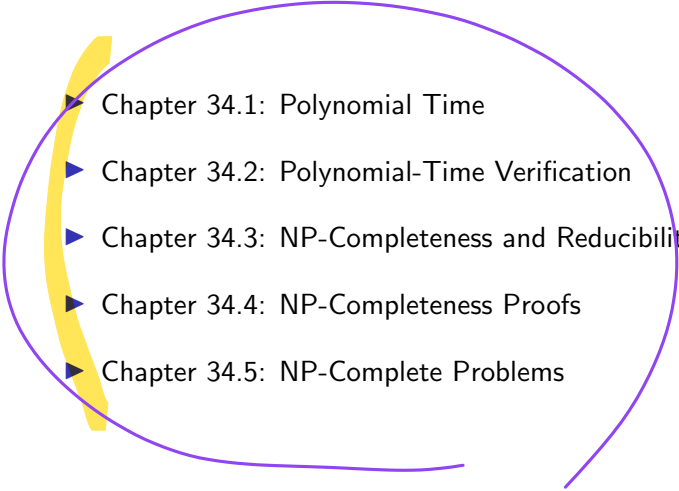
- ▶ $|V|$ calls to INSERT and EXTRACT-MIN
- ▶ At most $|E|$ calls to DECREASE-KEY

Total Time:

$$\mathcal{O}(|V| \log |V| + |E| \log |V|) = \mathcal{O}((|V| + |E|) \log |V|)$$

Use case: Efficient for sparse graphs

Chapter 34: NP-Completeness


- 
- ▶ Chapter 34.1: Polynomial Time
 - ▶ Chapter 34.2: Polynomial-Time Verification
 - ▶ Chapter 34.3: NP-Completeness and Reducibility
 - ▶ Chapter 34.4: NP-Completeness Proofs
 - ▶ Chapter 34.5: NP-Complete Problems

P, NP, and NP-Completeness: Informal Overview

- ▶ **P**: Problems solvable in polynomial time (e.g., $O(n^k)$ for some constant k).
- ▶ **NP**: Problems verifiable in polynomial time. Given a solution (*certificate*), it can be verified quickly.
- ▶ **NP-Complete (NPC)**: Problems in **NP** that are as hard as any other problem in NP.

If any NP-complete problem is in P, then $\overset{?}{\mathbf{P}} \overset{?}{=} \mathbf{NP}$.

Examples of NP Verifiability



NP-complete

Hamiltonian Cycle:

- ▶ Input: Graph $G = (V, E)$
- ▶ Certificate: Sequence $\langle v_1, v_2, \dots, v_{|V|} \rangle$
- ▶ Verify: Each node appears once; all edges exist; $v_{|V|} \rightarrow v_1$

P vs NP

- ▶ Every problem in **P** is in **NP**.
- ▶ Open question: Is **P** \subsetneq **NP**?
- ▶ If any NP-complete problem is in **P**, then **P** = **NP**.

?
 $P \neq NP$


Most researchers believe **P** \neq **NP**.

What is NP-Completeness?

A problem Q is NP-complete if:

1. $Q \in \mathbf{NP}$ (verifiable in poly time)
2. Every problem in \mathbf{NP} can be *reduced* to Q in poly time

If any NP-complete problem is solvable in poly time, all of NP is.



Why NP-Completeness Matters

- ▶ Proving NP-completeness shows likely intractability
- ▶ Better to design:
 - ▶ Approximation algorithms
 - ▶ Heuristics
 - ▶ Algorithms for special cases
- ▶ Many natural problems turn out to be NP-complete

Optimization vs Decision Problems

Optimization: Find best solution (e.g., shortest path, min cost).

Decision: Yes/No form. E.g.,

- ▶ **PATH:** Is there a path from u to v with $\leq k$ edges?

We study NP-completeness through decision problems.

Reductions: Relating Problem Hardness

- ▶ Given problems A and B
- ▶ If you can reduce A to B in poly time:
 - ▶ $A \leq_p B$: A is *no harder* than B
- ▶ Use known hard problems to show new problems are hard

Proving Intractability via Reductions

Suppose:

- ▶ Problem A is known to be hard (no poly algorithm)
- ▶ You reduce A to B (i.e., $A \leq_p B$)

Then:

- ▶ If B had a poly-time algorithm, so would A — contradiction!
- ▶ So B must also be hard

This is how we prove NP-completeness.

Chapter 34: NP-Completeness

- ▶ **Chapter 34.1: Polynomial Time**
- ▶ Chapter 34.2: Polynomial-Time Verification
- ▶ Chapter 34.3: NP-Completeness and Reducibility
- ▶ Chapter 34.4: NP-Completeness Proofs
- ▶ Chapter 34.5: NP-Complete Problems

Abstract Problems: Formal Definition

Abstract problem Q is defined as a **binary relation** between:

- ▶ A set I of **problem instances**, and
- ▶ A set S of **problem solutions**

Example: SHORTEST-PATH

- ▶ Instance: (G, u, v) — a graph and two vertices
- ▶ Solution: A path from u to v (or empty if none exists)
- ▶ Relation: Matches each instance with a valid shortest path

Note: There may be multiple valid solutions — shortest paths are not necessarily unique.

From Abstract Problems to Decision Problems

Why focus on decision problems?

- ▶ NP-completeness theory deals with problems having **yes/no answers**
- ▶ These are easier to define and analyze formally

A decision problem is a function:

$$f : I \rightarrow \{0, 1\}$$

- ▶ $f(i) = 1$ (yes) if the instance satisfies the condition
- ▶ $f(i) = 0$ (no) otherwise

Example: PATH

- ▶ Input: (G, u, v, k)
- ▶ Output: 1 if there exists a path from u to v with at most k edges; otherwise 0

Optimization vs. Decision Problems

Many problems are optimization problems:

- ▶ Find the shortest path, largest flow, minimum cut, etc.

But we can often transform them into decision problems:

- ▶ Ask: “*Is there a solution of cost $\leq k$?*”
- ▶ This transformation is key in NP-completeness proofs

Why this works:

- ▶ The decision version is *no harder* than the optimization version
- ▶ We can recover optimal values using repeated queries (e.g., binary search)

Optimization vs. Decision: Example

Problem: SHORTEST PATH

Optimization Version

- ▶ Given a graph G and vertices u, v
- ▶ **Goal:** Find the *length of the shortest path* from u to v

Decision Version

- ▶ Given a graph G , vertices u, v , and integer k
- ▶ **Question:** Is there a path from u to v with length $\leq k$?
- ▶ Answer is **Yes/No**

Both versions deal with the same underlying problem, but the decision version is used for NP-completeness analysis.

What is an Encoding?

Encoding is how we represent abstract problem instances as binary strings that a computer can understand.

Formal Definition: An *encoding* is a function:

$$e : S \rightarrow \{0, 1\}^*$$

that maps abstract objects S to binary strings.

Examples:

- ▶ $e(17) = 10001$
- ▶ $e(A) = 01000001$ (ASCII)
- ▶ A graph G : encoded as adjacency list, matrix, or edge list

Compound objects (sets, graphs, programs) are encoded by combining the encodings of their parts.

From Abstract to Concrete Problems

To run on a computer, an abstract decision problem must be encoded.

Concrete Problem:

- ▶ Input: a binary string $x \in \{0, 1\}^*$
- ▶ Output: $f(x) \in \{0, 1\}$

Size of instance: $n = |x|$ (length of binary string)

Polynomial-time algorithm: solves any input of length n in $O(n^k)$ for some constant k .

Encoding Can Affect Complexity

Encoding affects whether an algorithm appears to run in polynomial or exponential time.

Example: Integer k given as input

- ▶ Unary encoding: $111 \dots 1$ (length $n = k$) \rightarrow runtime $\Theta(k) = \Theta(n) \rightarrow$ polynomial
- ▶ Binary encoding: 10001 (length $n = \lfloor \log_2 k \rfloor + 1$) \rightarrow runtime $\Theta(k) = \Theta(2^n) \rightarrow$ exponential

Conclusion: The *choice of encoding* can drastically change the apparent complexity.

Polynomially Related Encodings

Goal: We want our notion of “efficient” (polynomial-time) to be encoding-independent.

Definition: Encodings e_1 and e_2 are *polynomially related* if:

- ▶ There exist functions f_{12}, f_{21} computable in polynomial time
- ▶ Such that $f_{12}(e_1(i)) = e_2(i)$ and $f_{21}(e_2(i)) = e_1(i)$ for all instances i

This ensures we can convert between encodings without changing the complexity class.

Lemma 34.1: Encoding-Invariant Polynomial Time

Lemma: Let Q be an abstract decision problem.
If encodings e_1 and e_2 are polynomially related, then:

$$e_1(Q) \in \mathbf{P} \iff e_2(Q) \in \mathbf{P}$$

Implication:

- ▶ Polynomial-time solvability doesn't depend on which *reasonable* encoding you choose
- ▶ We can now safely talk about abstract problems being in \mathbf{P} , assuming a standard encoding

Practical Takeaways on Encodings

- ▶ You must encode instances as binary strings to analyze complexity formally
- ▶ Bad encodings (e.g., unary) can make easy problems look hard
- ▶ As long as we use **reasonable encodings** (e.g., binary, base-3, ASCII), complexity classes are preserved
- ▶ We usually assume a *standard encoding* (e.g., $\langle G \rangle$ for a graph)

This allows us to focus on the *problem itself*, not the details of its representation.

A Formal-Language Framework for Decision Problems

Alphabet: A finite set of symbols (denoted Σ)

Language: A set of strings over Σ (i.e., $L \subseteq \Sigma^*$)

Examples:

- ▶ $\Sigma = \{0, 1\}$
- ▶ $L = \{10, 11, 101, 111, 10001, \dots\}$ (binary encodings of prime numbers)

Useful notation:

- ▶ ϵ : empty string
- ▶ Σ^* : set of all binary strings

Viewing Decision Problems as Languages

Decision problem $Q : \Sigma^* \rightarrow \{0, 1\}$ maps binary strings to yes/no.

We define the associated language:

$$L_Q = \{x \in \Sigma^* \mid Q(x) = 1\}$$

This is the set of inputs for which the answer is **yes**.

Example: PATH

- ▶ Input: $x = \langle G, u, v, k \rangle$ (encoded as a binary string)
- ▶ $x \in L_{\text{PATH}}$ if G contains a path from u to v of length $\leq k$

This lets us study decision problems as formal languages, enabling complexity class definitions.

Accepting vs. Deciding a Language

Algorithm A accepts x if $A(x) = 1$ (i.e., it outputs yes).

Language accepted by A :

$$L = \{x \in \Sigma^* \mid A(x) = 1\}$$

But: A might not halt on $x \notin L$ (it might loop forever).

Algorithm A decides L if:

- ▶ $A(x) = 1$ for $x \in L$
- ▶ $A(x) = 0$ for $x \notin L$

Polynomial-Time Acceptance vs. Decision

Accepted in polynomial time:

- ▶ A outputs 1 for $x \in L$ in $O(n^k)$ time
- ▶ No guarantee for $x \notin L$ (may loop forever)

Decided in polynomial time:

- ▶ A halts on all inputs
- ▶ $A(x) = 1$ if $x \in L$, and $A(x) = 0$ otherwise

Example: PATH can be both accepted and decided in polynomial time

Defining **P** via Formal Languages

Complexity Class **P**:

$\mathbf{P} = \{L \subseteq \Sigma^* \mid \text{there exists an algorithm that decides } L \text{ in } O(n^k) \text{ time}\}$

Theorem 34.2:

- ▶ **P** is the set of languages that are *accepted* by a polynomial-time algorithm
- ▶ Proof: Any polynomial-time accepting algorithm can be converted to a deciding algorithm via bounded simulation

Question?