

# Chapter 6. Heapsort

Joon Soo Yoo

April 3, 2025

# Assignment

- ▶ Read §6.3, §6.4, §6.5
- ▶ Problems
  - ▶ §6.1 - #2, 8
  - ▶ §6.2 - #2
  - ▶ §6.3 - #3, 4
  - ▶ §6.4 - #2, 4
  - ▶ §6.5 - #3, 7

## Chapter 6: Review

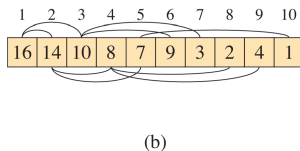
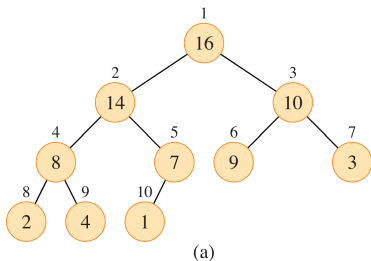
- ▶ **Chapter 6.1: Heaps**
- ▶ Chapter 6.2: Maintaining the heap property
- ▶ Chapter 6.3: Building a heap
- ▶ Chapter 6.4: The heapsort algorithm
- ▶ Chapter 6.5: Priority queues

# Heapsort: Overview

- ▶ Heapsort is a sorting algorithm introduced in this chapter.
- ▶ Like **Merge Sort**:
  - ▶ Runs in  $O(n \log n)$  time
- ▶ Like **Insertion Sort**:
  - ▶ Sorts **in place**, using only constant extra space
- ▶ Combines the best features of both: fast and space-efficient
- ▶ Introduces a new algorithmic idea: using a **data structure** called a **heap** to manage information

# What is a Binary Heap?

- ▶ A (binary) heap is a **nearly complete binary tree**
- ▶ Represented as an array  $A[1 : n]$
- ▶  $A.\text{heap\_size}$  denotes the number of elements currently in the heap
- ▶ Tree is filled top to bottom, left to right



# Heap Index Formulas

**Given a node at index  $i$ :**

$$\text{PARENT}(i) = \lfloor i/2 \rfloor$$

$$\text{LEFT}(i) = 2i$$

$$\text{RIGHT}(i) = 2i + 1$$

- ▶ These are computed efficiently using bit shifts
- ▶ Commonly implemented as inline functions or macros

# Heap Properties

- ▶ **Max-Heap:**

- ▶ For all nodes  $i$  (except the root), the following holds:

$$A[\text{parent}(i)] \geq A[i]$$

- ▶ So the largest element is at the root:  $A[1]$

- ▶ **Min-Heap:**

- ▶ For all nodes  $i$  (except the root), the following holds:

$$A[\text{parent}(i)] \leq A[i]$$

- ▶ So the smallest element is at the root:  $A[1]$

- ▶ Heapsort uses a **max-heap**

- ▶ Priority queues often use a **min-heap**

# Heap Height and Runtime

## Coming up in this chapter:

- ▶ MAX-HEAPIFY – maintain heap property ( $O(\log n)$ )
- ▶ BUILD-MAX-HEAP – make a heap from unordered array ( $\Theta(n)$ )
- ▶ HEAPSORT – sorting in-place using a heap ( $O(n \log n)$ )



# Chapter 6: Heapsort

- ▶ Chapter 6.1: Heaps
- ▶ **Chapter 6.2: Maintaining the heap property**
- ▶ Chapter 6.3: Building a heap
- ▶ Chapter 6.4: The heapsort algorithm
- ▶ Chapter 6.5: Priority queues

# Maintaining the Heap Property

- ▶ Core idea: MAX-HEAPIFY ensures the **max-heap property** is preserved.
- ▶ Input: Array  $A$  with  $A.\text{heap-size}$ , and index  $i$ .
- ▶ Assumes:
  - ▶ Subtrees rooted at  $\text{LEFT}(i)$  and  $\text{RIGHT}(i)$  are already max-heaps.
  - ▶ Node  $i$  may violate the heap property.
- ▶ Action: Let the value at  $A[i]$  **float down** the tree.

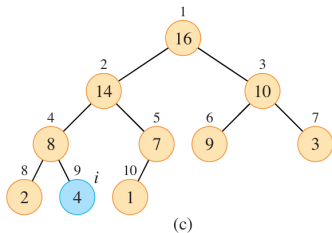
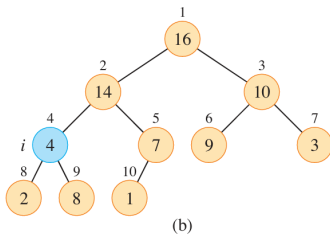
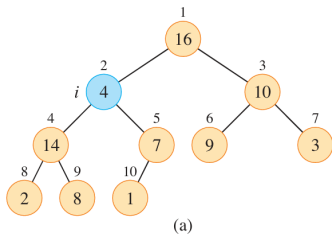
**Goal:** Make the subtree rooted at  $i$  a valid max-heap.

## MAX-HEAPIFY in Action (Figure 6.2)

- ▶ Step 1: Compare  $A[i]$ ,  $A[\text{LEFT}(i)]$ , and  $A[\text{RIGHT}(i)]$
- ▶ Step 2: Find the largest and set index in largest
- ▶ Step 3: If largest  $\neq i$ :
  - ▶ Swap  $A[i]$  and  $A[\text{largest}]$
  - ▶ Recursively call MAX-HEAPIFY on index largest

*Fixing violations from top-down.*

## MAX-HEAPIFY in Action (Figure 6.2)



*Fixing violations from top-down.*

# MAX-HEAPIFY Algorithm

---

**Algorithm 1** MAX-HEAPIFY( $A, i$ )

---

```
1:  $l \leftarrow \text{LEFT}(i)$ 
2:  $r \leftarrow \text{RIGHT}(i)$ 
3: if  $l \leq A.\text{heap-size}$  and  $A[l] > A[i]$  then
4:    $largest \leftarrow l$ 
5: else
6:    $largest \leftarrow i$ 
7: end if
8: if  $r \leq A.\text{heap-size}$  and  $A[r] > A[largest]$  then
9:    $largest \leftarrow r$ 
10: end if
11: if  $largest \neq i$  then
12:   exchange  $A[i]$  with  $A[largest]$ 
13:   MAX-HEAPIFY( $A, largest$ )
14: end if
```

---

**Note:**  $\text{LEFT}(i) = 2i$ ,  $\text{RIGHT}(i) = 2i + 1$

# Running Time of MAX-HEAPIFY

- ▶ Let  $T(n)$  be the worst-case time for a subtree with at most  $n$  nodes.
- ▶ Time to compare and swap at each step is  $O(1)$ .
- ▶ Recursive call occurs on a subtree of size at most  $\frac{2n}{3}$ .

$$T(n) \leq T\left(\frac{2n}{3}\right) + \Theta(1)$$

- ▶ Using the Master Theorem (Case 2):

$$T(n) = O(\log n)$$

**Alternatively:** Run time is  $O(h)$  where  $h$  is the height of node  $i$ .

# Think About It

## Question 1:

Why does each child subtree in MAX-HEAPIFY have size at most  $2n/3$ ?

**Hint:** Consider the structure of a complete binary tree.

## Question 2:

Given the recurrence:

$$T(n) \leq T(2n/3) + \Theta(1)$$

use the **Master Theorem** to find the time complexity of MAX-HEAPIFY.

## Answer: Why Subtree Size Is at Most $2n/3$

To analyze the worst-case size of a child subtree in a binary heap:

- ▶ Assume a heap with  $n$  nodes, rooted at index  $i$ .
- ▶ We want to find the maximum number of nodes in either the left or right subtree.
- ▶ The heap is a complete binary tree: all levels full, except possibly the last, which is filled left-to-right.
- ▶ To maximize the size of the left subtree, it should include as many nodes as possible in the last level.

Let  $k$  be the height of the last level, with:

$$\text{Left subtree size} = 2^k, \quad \text{Right subtree size} = 2^{k-1}$$

So:

$$2^k + 2^{k-1} = 3 \cdot 2^{k-1} \leq n \Rightarrow 2^{k-1} \leq \frac{n}{3} \Rightarrow \text{Max subtree size} \leq \frac{2n}{3}$$

**Conclusion:** Each recursive call to MAX-HEAPIFY acts on a subtree of size at most  $\boxed{2n/3}$ .



# Answer: Time Complexity via Master Theorem

We solve the recurrence:

$$T(n) \leq T(2n/3) + \Theta(1)$$

**Step 1:** Match to the Master Theorem form:

$$T(n) = T(n/b) + f(n)$$

Here,  $a = 1$ ,  $b = 3/2$  (since  $n/b = 2n/3$ ), and  $f(n) = \Theta(1)$

**Step 2:** Compare  $f(n)$  to  $n^{\log_b a} = n^{\log_{3/2} 1} = n^0 = 1$

**Conclusion:** This is Case 2 of the Master Theorem:

$$f(n) = \Theta(n^{\log_b a}) \Rightarrow T(n) = \Theta(\log n)$$

**Final Result:** MAX-HEAPIFY runs in  $O(\log n)$  time.

# Chapter 6: Heapsort

- ▶ Chapter 6.1: Heaps
- ▶ Chapter 6.2: Maintaining the heap property
- ▶ **Chapter 6.3: Building a heap**
- ▶ Chapter 6.4: The heapsort algorithm
- ▶ Chapter 6.5: Priority queues

## Exercise 6.1-8: Identifying Leaves in a Heap

**Statement:** In an  $n$ -element heap stored as an array  $A[1 : n]$ , the leaves are the nodes with indices:

$$\left\lfloor \frac{n}{2} \right\rfloor + 1, \left\lfloor \frac{n}{2} \right\rfloor + 2, \dots, n$$

**Proof:**

- ▶ In a heap, the children of a node  $i$  are at  $2i$  and  $2i + 1$ .
- ▶ If  $i > \left\lfloor \frac{n}{2} \right\rfloor$ , then  $2i > n \rightarrow i$  has no children  $\rightarrow$  it is a leaf.
- ▶ If  $i \leq \left\lfloor \frac{n}{2} \right\rfloor$ , then  $2i \leq n \rightarrow$  it has at least a left child  $\rightarrow$  it is internal.

**Conclusion:** Leaves are exactly those with indices from  $\left\lfloor \frac{n}{2} \right\rfloor + 1$  to  $n$ .

# BUILD-MAX-HEAP Algorithm

**Goal:** Convert an array  $A[1 : n]$  into a max-heap.

**Procedure:**

```
1: procedure BUILD-MAX-HEAP( $A, n$ )
2:    $A.\text{heap-size} \leftarrow n$ 
3:   for  $i \leftarrow \lfloor n/2 \rfloor$  downto 1 do
4:     MAX-HEAPIFY( $A, i$ )
5:   end for
6: end procedure
```

**Key Idea:** Internal nodes are heapified from bottom-up. Leaves are already heaps.

# Why BUILD-MAX-HEAP Works (Loop Invariant)

**Loop Invariant:** At the start of each iteration, nodes  $i + 1$  to  $n$  are all roots of valid max-heaps.

## Initialization:

- ▶  $i = \lfloor n/2 \rfloor$
- ▶ All nodes from  $i + 1$  to  $n$  are leaves  $\rightarrow$  trivially max-heaps.

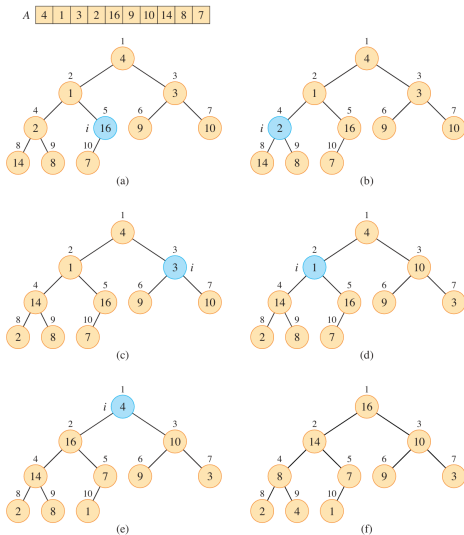
## Maintenance:

- ▶ Children of node  $i$  are greater than  $i$
- ▶  $\text{MAX-HEAPIFY}(A, i)$  ensures subtree rooted at  $i$  becomes a max-heap

## Termination:

- ▶ After loop, all nodes 1 through  $n$  are max-heap roots
- ▶ Especially  $A[1]$  — the root of the final heap

Figure: BUILD-MAX-HEAP in Action



# Naïve Runtime Analysis of BUILD-MAX-HEAP

- ▶ The procedure BUILD-MAX-HEAP(A) runs a loop:

for  $i = \lfloor n/2 \rfloor$  downto 1: MAX-HEAPIFY(A, i)

- ▶ Number of iterations:  $\lfloor n/2 \rfloor \approx O(n)$
- ▶ Each call to MAX-HEAPIFY can take up to  $O(\log n)$  time in the worst case (if the node is near the root)
- ▶ Therefore, total worst-case time is:

$$O(n) \cdot O(\log n) = O(n \log n)$$

- ▶ This is a correct upper bound — but not the tightest possible.

# Tighter Runtime Analysis of BUILD-MAX-HEAP

- ▶ The simple bound  $O(n \log n)$  assumes all calls to MAX-HEAPIFY take  $O(\log n)$  time
- ▶ In reality, many nodes are near the bottom of the heap and have small height
- ▶ MAX-HEAPIFY runs in  $O(h)$  where  $h$  is the height of the node
- ▶ Idea: sum the work done at each height
- ▶ Using known bounds on the number of nodes at each height:

$$\sum_{h=0}^{\lfloor \log n \rfloor} \left\lceil \frac{n}{2^{h+1}} \right\rceil \cdot O(h) = O(n)$$

- ▶ Tighter bound:  $O(n)$



# Key Questions for the Tighter Bound

To justify the tight  $O(n)$  analysis, we ask:

1. Why does an  $n$ -element heap have height  $\lfloor \log n \rfloor$ ?
2. Why are there at most  $\left\lceil \frac{n}{2^{h+1}} \right\rceil$  nodes of height  $h$ ?

These structural properties of binary heaps allow us to bound the total work done by BUILD-MAX-HEAP.

## Why is Heap Height $\lfloor \log_2 n \rfloor$ ?

**Goal:** Show that a binary heap with  $n$  elements has height  $\lfloor \log_2 n \rfloor$

- ▶ In a perfect binary tree of height  $h$ , number of nodes:

$$1 + 2 + 4 + \cdots + 2^h = 2^{h+1} - 1$$

- ▶ In a binary tree of height  $h$ , minimum number of nodes:

$$2^h \quad (\text{only one node at bottom level})$$

So the number of nodes  $n$  in a binary heap satisfies:

$$2^h \leq n < 2^{h+1}$$

Taking  $\log_2$ :

$$h \leq \log_2 n < h + 1 \Rightarrow h = \lfloor \log_2 n \rfloor$$

Why are there at most  $\left\lceil \frac{n}{2^{h+1}} \right\rceil$  nodes of height  $h$ ?

**Goal:** Bound the number of nodes at height  $h$  in a binary heap of size  $n$

- ▶ A node at height  $h$  must be the root of a subtree of height  $h$
- ▶ A full binary subtree of height  $h$  has:

$$\text{Total nodes} = 1 + 2 + 4 + \cdots + 2^h = 2^{h+1} - 1$$

- ▶ So each such node **uses up** at least  $2^{h+1} - 1$  positions in the heap
- ▶ Total number of such nodes is at most:

$$\left\lfloor \frac{n}{2^{h+1} - 1} \right\rfloor \leq \left\lceil \frac{n}{2^{h+1}} \right\rceil$$

# Tighter Runtime of BUILD-MAX-HEAP: Setup

**Let:**

- ▶  $c$ : constant factor in the  $O(h)$  cost of MAX-HEAPIFY
- ▶ Nodes at height  $h$ : at most  $\left\lceil \frac{n}{2^{h+1}} \right\rceil$

**Total cost:**

$$T(n) \leq \sum_{h=0}^{\lfloor \log n \rfloor} \left\lceil \frac{n}{2^{h+1}} \right\rceil \cdot ch$$

**Apply inequality:**  $\lceil x \rceil \leq 2x$  for all  $x \geq \frac{1}{2}$

$$\Rightarrow \left\lceil \frac{n}{2^{h+1}} \right\rceil \leq \frac{n}{2^h} \quad (\text{CLRS Exercise 6.3-2})$$

So:

$$T(n) \leq cn \sum_{h=0}^{\lfloor \log n \rfloor} \frac{h}{2^h}$$

# Transforming the Summation

We now simplify:

$$T(n) \leq cn \sum_{h=0}^{\lfloor \log n \rfloor} \frac{h}{2^h}$$

This summation is known to converge:

$$\sum_{h=0}^{\infty} \frac{h}{2^h} = 2$$

**From CLRS Equation A.11 (p.1142):**

$$\sum_{h=0}^{\infty} hx^h = \frac{x}{(1-x)^2} \quad \text{for } 0 < x < 1$$

$$\Rightarrow \sum_{h=0}^{\infty} \frac{h}{2^h} = \frac{\frac{1}{2}}{(1 - \frac{1}{2})^2} = \frac{1/2}{1/4} = 2$$

So we are safe to approximate:

$$T(n) \leq cn \cdot 2$$

# Final Result: BUILD-MAX-HEAP is Linear Time

**Final bound:**

$$T(n) \leq cn \cdot \sum_{h=0}^{\infty} \frac{1}{2^h} = cn \cdot 2 = O(n)$$

**Interpretation:**

- ▶ While a single MAX-HEAPIFY may cost  $O(\log n)$ , most nodes lie near the bottom of the heap
- ▶ Their small height makes their cost low
- ▶ So the total cost of building a heap is dominated by these shallow calls
- ▶ **Hence, BUILD-MAX-HEAP runs in linear time.**

|

# Chapter 6: Heapsort

- ▶ Chapter 6.1: Heaps
- ▶ Chapter 6.2: Maintaining the heap property
- ▶ Chapter 6.3: Building a heap
- ▶ **Chapter 6.4: The heapsort algorithm**
- ▶ Chapter 6.5: Priority queues

# Overview of HEAPSORT

## HEAPSORT Steps:

1. Call BUILD-MAX-HEAP to organize the input array into a max-heap.
2. The largest element is now at the root  $A[1]$ .
3. Swap  $A[1]$  with  $A[n]$  to move the max to its final sorted position.
4. Reduce the heap size by 1.
5. Call MAX-HEAPIFY on  $A[1]$  to restore the max-heap.
6. Repeat until heap size is 1.

Each iteration places the next-largest element in its final position.



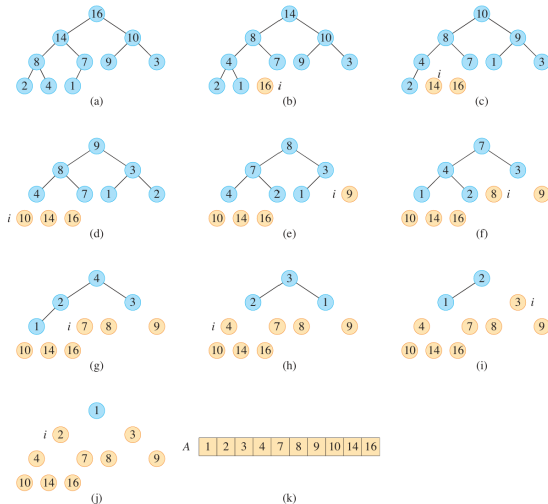
# HEAPSORT Pseudocode

HEAPSORT( $A, n$ ):

1. BUILD-MAX-HEAP( $A, n$ )
2. for  $i = n$  downto 2:
  - ▶ exchange  $A[1]$  with  $A[i]$
  - ▶  $A.\text{heap-size} = A.\text{heap-size} - 1$
  - ▶ MAX-HEAPIFY( $A, 1$ )

**Key Idea:** Each MAX-HEAPIFY fixes the heap with one fewer element.

# Visual Example of HEAPSORT



# Time Complexity of HEAPSORT

## Step-by-step costs:

- ▶ BUILD-MAX-HEAP takes  $O(n)$  time
- ▶ The loop runs  $n - 1$  times
- ▶ Each MAX-HEAPIFY call takes  $O(\log n)$

$$\Rightarrow T(n) = O(n) + (n - 1) \cdot O(\log n) = O(n \log n)$$

**HEAPSORT is an in-place, comparison-based sort with worst-case  $O(n \log n)$  time.**

# Chapter 6: Heapsort

- ▶ Chapter 6.1: Heaps
- ▶ Chapter 6.2: Maintaining the heap property
- ▶ Chapter 6.3: Building a heap
- ▶ Chapter 6.4: The heapsort algorithm
- ▶ **Chapter 6.5: Priority queues**

# What is a Priority Queue?

**Definition:** A **priority queue** is a data structure for maintaining a set  $S$  of elements, each with an associated value called a **key**.

## Operations Supported by a Max-Priority Queue:

- ▶  $\text{INSERT}(S, x, k)$ : Inserts element  $x$  with key  $k$  into  $S$
- ▶  $\text{MAXIMUM}(S)$ : Returns the element with the largest key
- ▶  $\text{EXTRACT-MAX}(S)$ : Removes and returns the element with the largest key
- ▶  $\text{INCREASE-KEY}(S, x, k)$ : Increases the key of  $x$  to new value  $k$ , where  $k$  is at least as large as the current key

# Application of Max-Priority Queues

## Use Case: Job Scheduling

- ▶ A computer shared by multiple users runs a scheduler.
- ▶ Jobs have priorities, and the scheduler selects the job with the highest priority.
- ▶ EXTRACT-MAX is used to select the next job to run.
- ▶ INSERT is used to add a new job at any time.

**Takeaway:** Max-priority queues are effective when the **most important task must be served first**.

# Application of Min-Priority Queues

## Min-Priority Queue Operations:

- ▶  $\text{INSERT}(S, x, k)$
- ▶  $\text{MINIMUM}(S)$
- ▶  $\text{EXTRACT-MIN}(S)$
- ▶  $\text{DECREASE-KEY}(S, x, k)$

## Use Case: Event-Driven Simulation

- ▶ Each event has a time of occurrence as its key.
- ▶ Events are simulated in chronological order.
- ▶  $\text{EXTRACT-MIN}$  selects the next event to simulate.
- ▶  $\text{INSERT}$  adds new future events during simulation.

**We will revisit min-priority queues in Chapters 21 and 22.**

# Objects and Keys in Priority Queues

## Recall:

- ▶ In sorting (e.g., heapsort), the array stores **keys** directly.
- ▶ Satellite data (extra information) is moved implicitly with keys.

## In Priority Queues:

- ▶ The heap array stores **pointers to objects**.
- ▶ Each object has a **key** attribute (e.g., `object.key`).
- ▶ Heap operations compare objects by their keys.

**Analogy:** Object = satellite data. Key = determines priority.



# MAX-HEAP-MAXIMUM and EXTRACT-MAX Code

## MAX-HEAP-MAXIMUM(A)

```
1: if A.heap-size < 1 then  
2:   error "heap underflow"  
3: end if  
4: return A[1]
```

## MAX-HEAP-EXTRACT-MAX(A)

```
1:  $max \leftarrow \text{MAX-HEAP-MAXIMUM}(A)$   
2:  $A[1] \leftarrow A[A.\text{heap-size}]$   
3:  $A.\text{heap-size} \leftarrow A.\text{heap-size} - 1$   
4:  $\text{MAX-HEAPIFY}(A, 1)$   
5: return  $max$ 
```

**Time Complexity:**  $\mathcal{O}(\log n)$

# MAXIMUM and EXTRACT-MAX Operations

## MAXIMUM:

- ▶ Returns the object with the largest key.
- ▶ Implemented as: `return A[1]`
- ▶ Runs in  $\mathcal{O}(1)$  time.

## EXTRACT-MAX:

- ▶ Removes and returns the maximum element.
- ▶ Replaces root with last element, reduces heap size, calls MAX-HEAPIFY.
- ▶ Runs in  $\mathcal{O}(\log n)$  time.
- ▶ Exchanges pointers (not raw data) and updates object-index mapping.

# MAX-HEAP-INCREASE-KEY Procedure

**Goal:** Increase the key of object  $x$  to a new value  $k$  in the heap  $A$ .

**Procedure:**

```
1: procedure MAX-HEAP-INCREASE-KEY( $A, x, k$ )
2:   if  $k < x.\text{key}$  then
3:     error "new key is smaller than current key"
4:   end if
5:    $x.\text{key} \leftarrow k$ 
6:   Find index  $i$  such that  $A[i] = x$ 
7:   while  $i > 1$  and  $A[\text{PARENT}(i)].\text{key} < A[i].\text{key}$  do
8:     Exchange  $A[i]$  with  $A[\text{PARENT}(i)]$ 
9:      $i \leftarrow \text{PARENT}(i)$ 
10:  end while
11: end procedure
```

**Time Complexity:**  $\mathcal{O}(\log n)$

# MAX-HEAP-INSERT Procedure

**Goal:** Insert object  $x$  into heap  $A$  with  $n$  slots.

**Procedure:**

```
1: procedure MAX-HEAP-INSERT( $A, x, n$ )
2:   if  $A$ .heap-size ==  $n$  then
3:     error "heap overflow"
4:   end if
5:    $A$ .heap-size  $\leftarrow A$ .heap-size + 1
6:    $k \leftarrow x$ .key
7:    $x$ .key  $\leftarrow -\infty$ 
8:    $A[A$ .heap-size]  $\leftarrow x$ 
9:   Map  $x$  to index  $A$ .heap-size
10:  MAX-HEAP-INCREASE-KEY( $A, x, k$ )
11: end procedure
```

**Time Complexity:**  $\mathcal{O}(\log n)$

# Summary: Max-Priority Queue with Max-Heap

**All operations are supported in  $\mathcal{O}(\log n)$  time:**

- ▶  $\text{MAXIMUM}(A) \rightarrow \mathcal{O}(1)$
- ▶  $\text{EXTRACT-MAX}(A) \rightarrow \mathcal{O}(\log n)$
- ▶  $\text{INCREASE-KEY}(A, x, k) \rightarrow \mathcal{O}(\log n)$
- ▶  $\text{INSERT}(A, x) \rightarrow \mathcal{O}(\log n)$

**Overhead:**

- ▶ Additional cost for maintaining object-to-index mapping (handles or hash tables)
- ▶ Typically  $\mathcal{O}(1)$  per update (expected)

# Question?