

Chapter 7. Quicksort

Joon Soo Yoo

April 9, 2025

Assignment

- ▶ Read §7
- ▶ Problems
 - ▶ §7.1 - #2, 4
 - ▶ §7.2 - #3, 6
 - ▶ §7.3 - #2
 - ▶ §7.4 - #1, 4

Chapter 7: Quicksort

- ▶ **Chapter 7.1: Description of quicksort**
- ▶ Chapter 7.2: Performance of quicksort
- ▶ Chapter 7.3: A randomized version of quicksort
- ▶ Chapter 7.4: Analysis of quicksort

Chapter 7: Quicksort — Introduction

- ▶ **Quicksort** is a divide-and-conquer sorting algorithm
- ▶ **Worst-case time:** $\Theta(n^2)$
- ▶ **Expected time (distinct elements):** $\Theta(n \log n)$
- ▶ Despite its worst case, it is often the **most practical sorting algorithm** in practice
- ▶ Advantages over merge sort:
 - ▶ **In-place** sorting (no extra memory)
 - ▶ Performs well even in virtual memory environments
 - ▶ Small hidden constants in the $\Theta(n \log n)$ bound

Quicksort is a Divide-and-Conquer Algorithm

Quicksort applies the **divide-and-conquer** paradigm introduced in Section 2.3.1:

- ▶ **Divide:** Partition the subarray $A[p \dots r]$ into two sides:
 - ▶ Elements $\leq A[q]$ go to the left (**low side**)
 - ▶ Elements $\geq A[q]$ go to the right (**high side**)
 - ▶ q is the final position of the pivot
- ▶ **Conquer:** Recursively apply quicksort to both subarrays
- ▶ **Combine:** Do nothing — the two sides are already sorted

Why No Combine Step in Quicksort?

Unlike Merge Sort, Quicksort has a trivial combine step.

- ▶ After partitioning:

$$A[p \dots q - 1] \leq A[q] \leq A[q + 1 \dots r]$$

- ▶ Recursive calls sort each side independently
- ▶ Since all elements on the left are $\leq A[q]$, and those on the right are $\geq A[q]$,

The full subarray $A[p \dots r]$ becomes sorted

No merging step is needed!

QUICKSORT(A, p, r) Pseudocode

```
QUICKSORT(A, p, r):  
    if p < r:  
        q = PARTITION(A, p, r)  
        // Recursively sort left and right parts  
        QUICKSORT(A, p, q - 1)  
        QUICKSORT(A, q + 1, r)
```

Initial call: QUICKSORT(A, 1, n) to sort the entire array.

Partitioning the Array

PARTITION(A, p, r) is the key subroutine in Quicksort.

- ▶ It rearranges the subarray $A[p \dots r]$ **in-place**
- ▶ The pivot is chosen as $x = A[r]$
- ▶ After execution:
 - ▶ All elements $\leq x$ appear to the left of the pivot
 - ▶ All elements $> x$ appear to the right of the pivot
 - ▶ The pivot is placed in its final sorted position
- ▶ Returns index q , the pivot's final position


PARTITION(A, p, r) Pseudocode

PARTITION(A, p, r)









```
1   $x = A[r]$  // the pivot
2   $i = p - 1$  // highest index into the low side
3  for  $j = p$  to  $r - 1$  // process each element other than the pivot
4      if  $A[j] \leq x$  // does this element belong on the low side?
5           $i = i + 1$  // index of a new slot in the low side
6          exchange  $A[i]$  with  $A[j]$  // put this element there
7  exchange  $A[i + 1]$  with  $A[r]$  // pivot goes just to the right of the low side
8  return  $i + 1$  // new index of the pivot
```

Partitioning: Four Regions

- ▶ The array $A[p \dots r]$ is divided into four regions during PARTITION:
 - ▶ **Brown** region $A[p \dots i]$: All values $\leq x$
 - ▶ **Blue** region $A[i + 1 \dots j - 1]$: All values $> x$
 - ▶ **White** region $A[j \dots r - 1]$: Values not yet examined
 - ▶ **Yellow** cell $A[r] = x$: The pivot
- ▶ These regions form the loop invariant for the for loop in PARTITION



Visualization



Loop Invariant for PARTITION

At the beginning of each iteration of the for loop:

1. If $p \leq k \leq i$, then $A[k] \leq x$ — the **tan region**
2. If $i + 1 \leq k \leq j - 1$, then $A[k] > x$ — the **blue region**
3. If $k = r$, then $A[k] = x$ — the **yellow region (pivot)**

Goal: Prove the loop invariant holds:

- ▶ Before the first iteration (Initialization)
- ▶ Maintained during each iteration (Maintenance)
- ▶ Guarantees correctness at the end (Termination)


Loop Invariant — Initialization

Before the first iteration:

- ▶ $i = p - 1, j = p$
- ▶ There are no elements between p and i , or between $i + 1$ and $j - 1$
- ▶ Therefore:
 - ▶ Conditions 1 and 2 of the invariant are trivially satisfied
 - ▶ Line 1 sets $x = A[r]$, so condition 3 is satisfied

The invariant holds before the loop starts.

Loop Invariant — Maintenance Case 1: $A[j] > x$




What happens: $A[j] > x$

- ▶ We do not increment i , no swap is performed
- ▶ We increment j

All loop invariant conditions continue to hold

- ▶ $A[j - 1]$ (just examined) now belongs in the blue region

Loop Invariant — Maintenance Case 2: $A[j] \leq x$



What happens: $A[j] \leq x$

- ▶ We increment i
- ▶ Swap $A[i]$ and $A[j]$
- ▶ Then increment j

Why it's correct:

- ▶ The new $A[i] \leq x \rightarrow$ condition 1 satisfied
- ▶ $A[j - 1] > x$ (from blue region) was just moved \rightarrow condition 2 holds

Loop Invariant — Termination

At termination:

- ▶ Loop runs for $r - p$ iterations \rightarrow ends with $j = r$
- ▶ The unexamined region $A[j \dots r - 1]$ is now empty
- ▶ All values are partitioned into three regions:
 - ▶ $A[p \dots i] \leq x$
 - ▶ $A[i + 1 \dots r - 1] > x$
 - ▶ $A[r] = x$

The final swap places x at position $i + 1$, and we return that index.

Loop invariant guarantees correctness of the partition.

Running Time of PARTITION

Claim: The running time of PARTITION on a subarray of size n is $\Theta(n)$

- ▶ The for loop runs from $j = p$ to $r - 1 \rightarrow$ exactly $n - 1$ iterations
- ▶ Each iteration performs a constant number of operations:
 - ▶ A comparison: $A[j] \leq x$
 - ▶ Maybe an increment of i
 - ▶ Maybe a swap between $A[i]$ and $A[j]$
- ▶ One final swap after the loop (pivot placement)

Conclusion: The total work is proportional to n , so the running time is:

$$\Theta(n)$$

Chapter 7: Quicksort

- ▶ Chapter 7.1: Description of quicksort
- ▶ **Chapter 7.2: Performance of quicksort**
- ▶ Chapter 7.3: A randomized version of quicksort
- ▶ Chapter 7.4: Analysis of quicksort

Quicksort Performance Depends on Pivot Balance

- ▶ The running time of Quicksort depends on how well the pivot splits the array.
- ▶ **Balanced partition:** Both sides of the partition are roughly equal in size
 - ▶ Quicksort behaves like Merge Sort: $\Theta(n \log n)$
- ▶ **Unbalanced partition:** One side is much larger than the other
 - ▶ Quicksort degrades to $\Theta(n^2)$, similar to Insertion Sort

Key Insight: The choice of pivot greatly affects performance

Best vs. Worst Case Intuition

Balanced Partitioning:

- ▶ Each pivot splits the array into two halves
- ▶ Recursion tree is roughly of height $\log n$
- ▶ Total time is $\Theta(n \log n)$

Unbalanced Partitioning:

- ▶ Each pivot gives a 1-element side and $n - 1$ element side
- ▶ Recursion tree becomes a long chain
- ▶ Total time is $\Theta(n^2)$

Quicksort Memory Usage

In-place sorting:

- ▶ Quicksort sorts the array in place (no extra arrays)
- ▶ This satisfies the definition of in-place sorting

But it still uses stack memory:

- ▶ Each recursive call uses stack space
- ▶ Stack depth = maximum depth of recursion tree
- ▶ In worst case (unbalanced), recursion depth = $\Theta(n)$
- ▶ In best case (balanced), depth = $\Theta(\log n)$

Worst-Case Partitioning in Quicksort

Worst case: Partitioning produces:

One subproblem of size $n - 1$, the other of size 0

- ▶ This occurs when the pivot is the smallest or largest element
- ▶ Happens consistently if input is already sorted (ascending or descending)

Time per partition: $\Theta(n)$

Recursive pattern:

$$T(n) = T(n - 1) + \Theta(n)$$

Solving the Worst-Case Recurrence

Given:

$$T(n) = T(n - 1) + \Theta(n)$$

Unrolling:

$$= T(n - 2) + \Theta(n - 1) + \Theta(n)$$

$$= \dots = T(1) + \sum_{k=2}^n \Theta(k) = \Theta(n^2)$$

Result: Worst-case time is:

$$T(n) = \Theta(n^2)$$

Best-Case Partitioning in Quicksort

Best case: Every call to PARTITION splits the array into two equal halves.

- ▶ One subarray of size $\lfloor \frac{n-1}{2} \rfloor \leq \frac{n}{2}$
- ▶ One subarray of size $\lceil \frac{n-1}{2} \rceil \leq \frac{n}{2}$
- ▶ Partitioning work per level: $\Theta(n)$

Running time recurrence:

$$T(n) = 2T(n/2) + \Theta(n)$$

Apply Master Theorem (Case 2):

$$T(n) = \Theta(n \log n)$$

Conclusion: When Quicksort always partitions evenly, its runtime matches Merge Sort: $\Theta(n \log n)$

Balanced Partitioning in Quicksort

Idea: Even when partitioning is *not perfectly even*, Quicksort performs well.


- ▶ Example: PARTITION always splits the array in a 9-to-1 ratio
- ▶ This seems unbalanced — but Quicksort still achieves:

$$T(n) = T(9n/10) + T(n/10) + \Theta(n)$$

- ▶ The cost at each level is $\Theta(n)$

Goal: Understand how such a split leads to $O(n \log n)$ time

Recursion Tree for a 9-to-1 Partitioning



- ▶ Each level of the recursion tree does $O(n)$ total work
- ▶ The depth of the tree is:

$$\Theta(\log_{10/9} n) = \Theta(\log n)$$

- ▶ Total work:

$$O(n) \cdot \Theta(\log n) = O(n \log n)$$

Constant Proportional Splits Yield $O(n \log n)$

- ▶ Even a 99-to-1 partition gives $O(n \log n)$ time
- ▶ As long as each partition is reduced by a constant factor at every level
- ▶ The key: tree depth is still $\Theta(\log n)$, and each level does $O(n)$ work

Therefore:

$$T(n) = O(n \log n) \quad \text{for any constant-ratio split}$$

The split ratio affects only the constant factor, not the asymptotic growth



Intuition for the Average Case

Goal: Understand why Quicksort is fast *on average*, despite some bad splits.

- ▶ Input is modeled as a random permutation of distinct elements
- ▶ Quicksort's performance depends on the relative ordering, not values
- ▶ Like the hiring problem (Chapter 5), we assume all orderings are equally likely

Key idea: Average-case behavior is a mix of good and bad splits, but still leads to efficient sorting.

Alternating Splits



- ▶ First level: worst-case split \rightarrow sizes $n - 1$ and 0
- ▶ Second level: good split \rightarrow halves of $n - 1$
- ▶ Total work for two levels:

$$\Theta(n) + \Theta(n - 1) = \Theta(n)$$

- ▶ Even though the first split was bad, the good split catches up

Conclusion: Average Case is $O(n \log n)$

- ▶ In the average case, good and bad splits are **randomly mixed**
- ▶ Even alternating between worst-case and best-case still gives:

$$\Theta(n \log n)$$

- ▶ Bad splits get “absorbed” by good ones
- ▶ So average-case cost is close to best case — with only a slightly larger constant

Result: Quicksort is fast on average, even with occasional bad pivots!

Chapter 7: Quicksort

- ▶ Chapter 7.1: Description of quicksort
- ▶ Chapter 7.2: Performance of quicksort
- ▶ **Chapter 7.3: A randomized version of quicksort**
- ▶ Chapter 7.4: Analysis of quicksort

Why Randomize Quicksort?

Problem: Our average-case analysis assumes the input is a random permutation

- ▶ In real-world scenarios, input may be sorted, reverse-sorted, or adversarial
- ▶ In such cases, deterministic Quicksort can degrade to $\Theta(n^2)$
- ▶ **Solution:** Use randomization to guard against worst-case inputs

Benefit: With random pivot selection, all inputs behave like "random" inputs on average

How Randomized Quicksort Works

Idea: Instead of always choosing the last element $A[r]$ as the pivot:

- ▶ Choose a random index $i \in [p, r]$
- ▶ Swap $A[i]$ with $A[r]$
- ▶ Then proceed with normal partitioning

Effect: Every element in $A[p \dots r]$ has equal probability of being the pivot

RANDOMIZED-QUICKSORT and PARTITION

RANDOMIZED-PARTITION(A, p, r)

```
1   $i = \text{RANDOM}(p, r)$   
2  exchange  $A[r]$  with  $A[i]$   
3  return PARTITION( $A, p, r$ )
```

RANDOMIZED-QUICKSORT(A, p, r)

```
1  if  $p < r$   
2       $q = \text{RANDOMIZED-PARTITION}(A, p, r)$   
3      RANDOMIZED-QUICKSORT( $A, p, q - 1$ )  
4      RANDOMIZED-QUICKSORT( $A, q + 1, r$ )
```

Only change: one random swap before partitioning

Why Randomization Helps

- ▶ Avoids dependence on input order
- ▶ Every pivot has equal probability \rightarrow partitions are likely balanced on average
- ▶ All inputs behave like random permutations
- ▶ Prevents adversarial inputs from triggering worst-case behavior

Result: Expected runtime becomes $O(n \log n)$ for *all* inputs

Chapter 7: Quicksort

- ▶ Chapter 7.1: Description of quicksort
- ▶ Chapter 7.2: Performance of quicksort
- ▶ Chapter 7.3: A randomized version of quicksort
- ▶ **Chapter 7.4: Analysis of quicksort**

Formal Analysis of Quicksort

Goal: Prove that the worst-case running time of Quicksort is $\Theta(n^2)$

- ▶ Section 7.2 gave an intuition for how bad partitioning leads to $\Theta(n^2)$
- ▶ Now we'll use the **substitution method** to formally prove:

$$T(n) = O(n^2)$$

- ▶ This worst-case applies to both QUICKSORT and RANDOMIZED-QUICKSORT

Worst-Case Recurrence for Quicksort

Let $T(n)$ be the worst-case time for QUICKSORT on input of size n .

In the worst case:

- ▶ The pivot causes a maximally unbalanced split: $q = 0$ or $q = n - 1$
- ▶ PARTITION does $\Theta(n)$ work

Worst-case recurrence:

$$T(n) = \max_{0 \leq q \leq n-1} \{T(q) + T(n-1-q)\} + \Theta(n) \quad (\text{Equation 7.1})$$

Using the Substitution Method

Goal: Prove $T(n) = O(n^2)$ using substitution.

Step 1: Guess the bound (inductive hypothesis)

Assume:

$$T(m) \leq cm^2 \quad \text{for all } m < n, \text{ where } c > 0$$

Step 2: Apply the recurrence to $T(n)$

From the recurrence:

$$T(n) = \max_{0 \leq q \leq n-1} \{T(q) + T(n-1-q)\} + \Theta(n)$$

Apply the inductive hypothesis to the recursive calls:

$$T(n) \leq \max_{0 \leq q \leq n-1} \{cq^2 + c(n-1-q)^2\} + \Theta(n)$$

Now simplify and bound the maximum to complete the proof.

Bounding the Maximum

Focus on the inner term:

$$q^2 + (n - 1 - q)^2$$

Expand:

$$= q^2 + (n - 1)^2 - 2q(n - 1) + q^2 = (n - 1)^2 + 2q(q - (n - 1))$$

The second term is always ≤ 0 , so:

$$q^2 + (n - 1 - q)^2 \leq (n - 1)^2$$

So:

$$T(n) \leq c(n - 1)^2 + \Theta(n) \leq cn^2 - c(2n - 1) + \Theta(n) \leq cn^2$$

(for large enough c)

Conclusion: Worst-Case is $\Theta(n^2)$

- ▶ We showed $T(n) = O(n^2)$ via substitution method
- ▶ We also know that $T(n) = \Omega(n^2)$ in the worst-case input (exercise 7.4-1)

Therefore:

$$T(n) = \Theta(n^2)$$

Running Time and Comparisons in Quicksort

```
PARTITION( $A, p, r$ )  
1   $x = A[r]$  // the pivot  
2   $i = p - 1$  // highest index into the low side  
3  for  $j = p$  to  $r - 1$  // process each element other than the pivot  
4      if  $A[j] \leq x$  // does this element belong on the low side?  
5           $i = i + 1$  // index of a new slot in the low side  
6          exchange  $A[i]$  with  $A[j]$  // put this element there  
7  exchange  $A[i + 1]$  with  $A[r]$  // pivot goes just to the right of the low side  
8  return  $i + 1$  // new index of the pivot
```

Observation: Quicksort's total running time is primarily due to comparisons in the PARTITION procedure.

Focus: Count comparisons of elements (not indices)

- ▶ Occur in line 4 of PARTITION
- ▶ Each compares an element to the current pivot

Let: X = Total number of element comparisons

Then the total running time is tied to X

Lemma 7.1 – Running Time in Terms of Comparisons

Lemma 7.1: The running time of QUICKSORT on an n -element array is:

$$O(n + X)$$

where X is the number of element comparisons made in line 4 of PARTITION

This lemma applies to both:

- ▶ QUICKSORT (deterministic)
- ▶ RANDOMIZED-QUICKSORT

Proof Idea of Lemma 7.1

- ▶ Each call to PARTITION removes the pivot from future calls
- ▶ There are at most n calls to PARTITION (one per pivot)
- ▶ Each QUICKSORT call can make up to 2 recursive calls \rightarrow at most $2n$ total calls

For each PARTITION call:

- ▶ Outside the loop: $O(1)$
- ▶ Loop iterations: one comparison per iteration

Total time:

$$O(n) \text{ (overhead)} + O(X) \text{ (comparisons)} = O(n + X)$$

Goal: Expected Number of Comparisons in Quicksort

Let X = total number of element comparisons made by PARTITION

Goal: Compute $\mathbb{E}[X]$

- ▶ Each comparison happens in line 4 of PARTITION
- ▶ To understand $\mathbb{E}[X]$, we must ask:
 - ▶ When are two elements compared?
 - ▶ How many such comparisons occur?

Tool: Analyze using the order of elements in the sorted array

$$z_1 < z_2 < \cdots < z_n$$

Lemma 7.2 – When Are Two Elements Compared?

Let $z_i < z_j$ be two elements in the sorted array. Then:

Lemma 7.2:

z_i and z_j are compared if and only if

either z_i or z_j is the first pivot chosen from $Z_{ij} = \{z_i, \dots, z_j\}$

Consequences:

- ▶ If any z_k with $i < k < j$ is chosen first $\rightarrow z_i$ and z_j fall into different sides \rightarrow no comparison
- ▶ Comparison happens only once — when the pivot is chosen

No Two Elements Are Compared Twice

Lemma 7.2 (continued): Once two elements z_i and z_j are compared, they are never compared again.

Why?

- ▶ The comparison happens only if one of z_i or z_j is chosen first as pivot in Z_{ij}
- ▶ After that pivot is chosen, it is removed from the array
- ▶ So z_i and z_j can never appear together in the same recursive call again

Example: Comparing Elements in Randomized Quicksort

Suppose the input is $\{1, 2, \dots, 10\}$ (random order)

- ▶ First pivot chosen is 7 \rightarrow PARTITION splits array into:

$$\{1, 2, 3, 4, 5, 6\} \quad \text{and} \quad \{8, 9, 10\}$$

- ▶ 7 is compared with every element — it's the first pivot in each Z_{ij}
- ▶ 2 and 9 will **never be compared** because 7 splits them
- ▶ 7 and 9 **are compared** because 7 is the first pivot in $Z_{7,9}$

Takeaway: A pair (z_i, z_j) is only compared if no element in between was chosen as pivot first.

Lemma 7.3 – Probability of Comparing Two Elements

Let $z_1 < z_2 < \dots < z_n$ be distinct elements in sorted order.

Lemma 7.3: For any $i < j$, the probability that z_i and z_j are compared during Randomized Quicksort is:

$$\Pr[z_i \text{ compared with } z_j] = \frac{2}{j - i + 1}$$

Why Are z_i and z_j Compared?

- ▶ Let $Z_{ij} = \{z_i, z_{i+1}, \dots, z_j\}$
- ▶ z_i and z_j are compared **iff** one of them is chosen first as pivot from Z_{ij}
- ▶ Each element in Z_{ij} has equal chance to be the first pivot

$$\Pr[\text{first pivot is } z_i] = \Pr[\text{first pivot is } z_j] = \frac{1}{j - i + 1}$$

$$\Pr[z_i \text{ and } z_j \text{ are compared}] = \frac{1}{j - i + 1} + \frac{1}{j - i + 1} = \boxed{\frac{2}{j - i + 1}}$$

Theorem 7.4 – Expected Time of Randomized Quicksort

Theorem 7.4: The expected running time of RANDOMIZED-QUICKSORT on n distinct elements is:

$$O(n \log n)$$

Strategy: We'll compute the expected number of comparisons $\mathbb{E}[X]$, and use:

$$\text{Total time} = O(n + X) \quad (\text{Lemma 7.1})$$

Step 1: Define Indicator Random Variables

Let elements be sorted: $z_1 < z_2 < \dots < z_n$

Define:

$$X_{ij} = \begin{cases} 1 & \text{if } z_i \text{ is compared with } z_j \\ 0 & \text{otherwise} \end{cases} \quad \text{for } 1 \leq i < j \leq n$$

Then the total number of comparisons:

$$X = \sum_{1 \leq i < j \leq n} X_{ij}$$

We will compute $\mathbb{E}[X]$ using linearity of expectation.

Step 2: Apply Linearity of Expectation

$$\mathbb{E}[X] = \mathbb{E}\left[\sum_{1 \leq i < j \leq n} X_{ij}\right] = \sum_{1 \leq i < j \leq n} \mathbb{E}[X_{ij}]$$

From Lemma 5.1 (Chapter 5), $\mathbb{E}[X_{ij}] = \Pr[z_i \text{ compared with } z_j]$
And from Lemma 7.3:

$$\mathbb{E}[X_{ij}] = \frac{2}{j - i + 1}$$

So:

$$\mathbb{E}[X] = \sum_{1 \leq i < j \leq n} \frac{2}{j - i + 1}$$

Step 3: Change of Variables to Simplify the Sum

Change variable: let $k = j - i$, so:

$$\mathbb{E}[X] = \sum_{i=1}^{n-1} \sum_{k=1}^{n-i} \frac{2}{k+1} < \sum_{i=1}^{n-1} \sum_{k=1}^n \frac{2}{k}$$

The inner sum is a harmonic series:

$$\sum_{k=1}^n \frac{1}{k} = O(\log n)$$

So the total becomes:

$$\mathbb{E}[X] = O(n \log n)$$

Conclusion: Expected Running Time

We proved:

$$\mathbb{E}[X] = O(n \log n) \quad (\text{expected number of comparisons})$$

And from Lemma 7.1:

$$\text{Expected running time of Quicksort} = O(n + X) = O(n \log n)$$

Quicksort is fast on average for any input!

Question?