

Chapter 8. Sorting in Linear Time

Joon Soo Yoo

April 9, 2025

Assignment

- ▶ Read §8
- ▶ Problems
 - ▶ §8.2 - 3
 - ▶ §8.3 - 3, 5
 - ▶ §8.4 - 2


Chapter 8: Sorting in Linear Time

- ▶ **Chapter 8.1: Lower Bounds for Sorting**
- ▶ Chapter 8.2: Counting Sort
- ▶ Chapter 8.3: Radix Sort
- ▶ Chapter 8.4: Bucket Sort

Why Assume Distinct Elements in Lower Bound Proofs?

- ▶ We aim to prove a **lower bound** for comparison-based sorting algorithms.
- ▶ Sorting with **duplicates** can be easier:
 - ▶ Some comparisons may be skipped.
 - ▶ Fewer rearrangements may be required.
- ▶ To make the argument general, we assume all elements are **distinct**.
- ▶ **Key Insight:** Proving a lower bound for distinct elements means the bound also holds when duplicates are allowed.
- ▶ Comparisons like $a_i = a_j$ are useless when all elements are distinct.
- ▶ So, we only consider comparisons of the form $a_i < a_j$.

The Decision-Tree Model



- ▶ A **comparison sort** uses only element comparisons like $a_i < a_j$ to determine order.
- ▶ We model these sorts using a **decision tree**:
 - ▶ A full binary tree where each internal node represents a comparison $a_i : a_j$
 - ▶ Each leaf corresponds to a permutation of the input
- ▶ The path from root to leaf = sequence of comparisons made for a specific input

Key Properties of the Decision Tree

- ▶ Internal nodes: comparisons $a_i : a_j$
- ▶ Leaves: permutations $\langle \pi(1), \pi(2), \dots, \pi(n) \rangle$
- ▶ The tree must have at least $n!$ **reachable leaves** to represent all permutations
- ▶ Execution follows a single root-to-leaf path, depending on comparison outcomes
- ▶ The height of the tree = **worst-case number of comparisons**

Theorem 8.1: Lower Bound for Comparison Sorts

Statement: Any comparison sort algorithm requires $\Omega(n \log n)$ comparisons in the worst case.

Proof:

- ▶ A correct comparison sort must distinguish all $n!$ input permutations.
- ▶ Therefore, its decision tree has at least $n!$ **reachable leaves**.
- ▶ Any binary tree of height h has at most 2^h leaves.

$$n! \leq 2^h \Rightarrow h \geq \log_2(n!)$$

- ▶ Apply Stirling's approximation:

$$n! \approx \sqrt{2\pi n} \left(\frac{n}{e}\right)^n \Rightarrow \log_2(n!) = \Theta(n \log n)$$

- ▶ Thus, worst-case comparisons: $\Omega(n \log n)$

Corollary 8.2: Asymptotic Optimality of Heapsort and Merge Sort

Corollary: *Heapsort and merge sort are asymptotically optimal comparison sorts.*

Justification:

- ▶ By Theorem 8.1, any comparison sort requires $\Omega(n \log n)$ comparisons in the worst case.
- ▶ Both heapsort and merge sort achieve a worst-case time of $O(n \log n)$.
- ▶ Therefore, they match the lower bound up to constant factors:

$$\text{Worst-case time} = \Theta(n \log n)$$

Conclusion: No comparison sort is asymptotically faster than heapsort or merge sort.

Chapter 8: Sorting in Linear Time

- ▶ Chapter 8.1: Lower Bounds for Sorting
- ▶ **Chapter 8.2: Counting Sort**
- ▶ Chapter 8.3: Radix Sort
- ▶ Chapter 8.4: Bucket Sort

Counting Sort — Introduction

- ▶ Counting sort assumes that the input array contains n integers in the range 0 to k , for some integer k .
- ▶ It runs in $\Theta(n + k)$ time.
- ▶ **When $k = O(n)$, counting sort runs in $\Theta(n)$ time.**
- ▶ Unlike comparison sorts, it uses the **value of the input** to determine position.
- ▶ Key idea:
 - ▶ For each input x , count how many elements are $\leq x$
 - ▶ Place x directly into its final position in the output array

Example: If 17 elements are $\leq x$, then x goes into position 17.

Understanding COUNTING-SORT Step-by-Step

1. **Initialize:** Zero out array $C[0 \dots k]$
2. **Count:** For each value x in A , increment $C[x]$
3. **Accumulate:** Compute how many values are $\leq i$ by cumulative sum
4. **Place:** For each $A[j]$, set $B[C[A[j]]] = A[j]$, and then decrement $C[A[j]]$

Key point: Decrementing $C[A[j]]$ handles duplicates by placing them in stable, ordered positions.


Note: We process A in reverse (from n to 1) to ensure **stability**.

COUNTING-SORT Procedure


COUNTING-SORT(A, n, k)

```
1  let  $B[1 : n]$  and  $C[0 : k]$  be new arrays
2  for  $i = 0$  to  $k$ 
3       $C[i] = 0$ 
4  for  $j = 1$  to  $n$ 
5       $C[A[j]] = C[A[j]] + 1$ 
6  //  $C[i]$  now contains the number of elements equal to  $i$ .
7  for  $i = 1$  to  $k$ 
8       $C[i] = C[i] + C[i - 1]$ 
9  //  $C[i]$  now contains the number of elements less than or equal to  $i$ .
10 // Copy  $A$  to  $B$ , starting from the end of  $A$ .
11 for  $j = n$  downto 1
12      $B[C[A[j]]] = A[j]$ 
13      $C[A[j]] = C[A[j]] - 1$  // to handle duplicate values
14 return  $B$ 
```


COUNTING-SORT Diagram




(a)




(b)




(c)



(d)



(e)



(f)

Time Complexity of Counting Sort

- ▶ Let n : number of input elements, k : maximum value in the input
- ▶ Loop breakdown:
 - ▶ Lines 2–3: Initialize $C[0 \dots k] \rightarrow \Theta(k)$
 - ▶ Lines 4–5: Count frequencies from $A \rightarrow \Theta(n)$
 - ▶ Lines 7–8: Compute prefix sums $\rightarrow \Theta(k)$
 - ▶ Lines 11–13: Place elements into $B \rightarrow \Theta(n)$
- ▶ **Total running time:**

$$\Theta(n + k)$$

- ▶ When $k = O(n)$, counting sort runs in $\Theta(n)$ time.

Why Counting Sort Can Beat $\Omega(n \log n)$

- ▶ In Section 8.1, we proved that any comparison sort must take $\Omega(n \log n)$ in the worst case.
- ▶ **Counting sort is not a comparison sort:**
 - ▶ It performs no comparisons between input elements.
 - ▶ Instead, it uses the **actual values** as indices into array C .
- ▶ Therefore, the lower bound of $\Omega(n \log n)$ does not apply.

Stability of Counting Sort

- ▶ **Definition:** A sorting algorithm is **stable** if it preserves the relative order of equal elements.
- ▶ In counting sort, if two elements have the same value, the one appearing first in the input appears first in the output.
- ▶ Stability matters when elements carry **satellite data**.
- ▶ **Example:**
 - ▶ Input: [(5, A), (3, B), (5, C)]
 - ▶ Stable output: [(3, B), (5, A), (5, C)]
- ▶ **Conclusion:** Counting sort's stability is **required** for radix sort to work correctly.


Chapter 8: Sorting in Linear Time

- ▶ Chapter 8.1: Lower Bounds for Sorting
- ▶ Chapter 8.2: Counting Sort
- ▶ **Chapter 8.3: Radix Sort**
- ▶ Chapter 8.4: Bucket Sort

Radix Sort and Card-Sorting Machines

- ▶ Radix sort was originally implemented using **mechanical card-sorting machines**.
- ▶ Each card had **80 columns**.
- ▶ In each column, a hole could be punched in **1 of 12 positions**.
- ▶ The sorter could be mechanically programmed to:
 - ▶ Examine a specific column
 - ▶ Distribute each card into one of 12 bins based on the punched position
- ▶ An operator would then **gather cards bin by bin**, so that cards with the first punch appear on top.

Understanding Punch Card Encoding (IBM Standard)



- ▶ The top two unlabeled rows correspond to **zone punches**:
 - ▶ Top row = **row 12**
 - ▶ Second row = **row 11**
- ▶ Digits 0–9 are encoded by a single punch in rows 0–9.
- ▶ Letters are encoded by combining for instance, **Row 12 + 1–9 = A–I**

Encoding Digits in Punch Cards

- ▶ For decimal digits, each column used only **10 of the 12 punch positions**:
 - ▶ Rows 0 through 9 \rightarrow digits 0 through 9
 - ▶ Rows 11 and 12 reserved for encoding letters or special characters
- ▶ A d -digit number occupies a field of **d columns**.
- ▶ Since the card sorter can look at only **one column at a time**, sorting n cards by a d -digit number requires a sorting algorithm that processes each digit.

The Need for a Sorting Algorithm

- ▶ You might intuitively think to sort numbers by the **most significant digit (MSD)**:
 - ▶ Sort by the leftmost digit
 - ▶ Recursively sort each bin by the next digit
 - ▶ Combine the bins in order to complete the sort
- ▶ This seems natural, but in physical systems, it presents a practical challenge...

Drawback of MSD Sorting


- ▶ MSD-based sorting creates **many intermediate piles** of cards.
- ▶ After sorting by the first digit, you have 10 bins.
- ▶ To recursively sort one bin, the other 9 bins must be **set aside and tracked**.
- ▶ This results in increased space and logistical complexity, especially in manual sorting.

Conclusion: MSD sorting leads to extra work managing intermediate piles—hence a better algorithmic strategy is needed.

Radix Sort: LSD-First Strategy and Stability

- ▶ Radix sort solves the problem of card sorting **counterintuitively** by sorting on the **least significant digit first**.
- ▶ Cards are first sorted by the rightmost digit:
 - ▶ Cards go into bins (0–9) based on that digit.
 - ▶ Bins are gathered in order: bin 0, then bin 1, ..., bin 9.
- ▶ Then the entire deck is sorted again on the next digit to the left, and recombined in the same way.
- ▶ This process continues until all d digits have been sorted.
- ▶ **Remarkably**, at that point the cards are fully sorted by their entire d -digit number.
- ▶ Thus, radix sort requires only d passes through the deck.

Stability in Radix Sort



- ▶ For radix sort to work correctly, the digit-level sorting must be **stable**.
- ▶ Stability ensures that the relative order of elements with equal digits is preserved.
- ▶ Mechanical card sorters are stable by design.

Radix Sort for Multi-Field Sorting (e.g., Dates)

- ▶ Radix sort can be used to sort **records with multiple fields**.
- ▶ Example: sorting dates by **year**, **month**, and **day**.
- ▶ One approach:
 - ▶ Use a comparison sort with a comparator that checks year first, then month, then day.
- ▶ Alternatively, use radix sort:
 - ▶ **First pass:** sort by day (least significant)
 - ▶ **Second pass:** sort by month
 - ▶ **Final pass:** sort by year (most significant)
- ▶ Requires each sort to be **stable**.

RADIX-SORT Pseudocode

RADIX-SORT assumes:

- ▶ Each element in array $A[1 \dots n]$ has d digits
- ▶ Digit 1 is the **least significant digit (LSD)**, and digit d is the most significant

RADIX-SORT(A, n, d)

```
1  for  $i = 1$  to  $d$ 
2      use a stable sort to sort array  $A[1 : n]$  on digit  $i$ 
```

- ▶ The pseudocode does not specify which stable sort to use
- ▶ **COUNTING-SORT** is commonly used for digit-level sorting
- ▶ Why counting sort?
 - ▶ Runs in $\Theta(n + k)$, where k is the digit range (e.g., 0–9)
 - ▶ Naturally stable and efficient for small-digit alphabets

Time Complexity of RADIX-SORT

Lemma 8.3.

Given n d -digit numbers in which each digit can take on up to k possible values, **RADIX-SORT correctly sorts** these numbers in

$$\Theta(d(n + k))$$

time, assuming the stable sort used runs in $\Theta(n + k)$.

Proof:

- ▶ Each pass processes n items, each with digit range 0 to $k - 1$.
- ▶ COUNTING-SORT takes $\Theta(n + k)$ per pass.
- ▶ With d passes, total running time is:

$$\Theta(d(n + k))$$

Corollary: Linear-Time Sorting When d and $k = O(n)$

- ▶ If d is constant and $k = O(n)$, then:

$$\Theta(d(n + k)) = \Theta(n)$$

- ▶ So **radix sort runs in linear time** in this case.
- ▶ We also have flexibility in how we define “digits,” especially for binary keys.

Lemma 8.4: Binary Keys in Radix Sort

Lemma 8.4.

Given n binary keys of b bits each, and any positive integer $r \leq b$, radix sort can sort the keys in:

$$\Theta\left(\frac{b}{r}(n + 2^r)\right)$$

time, assuming the stable sort used takes $\Theta(n + k)$ time for inputs in the range $[0, k]$.

Proof of Lemma 8.4

- ▶ A binary key is just a number made of b bits (e.g., 32-bit integers).
- ▶ Radix sort requires us to define “digits” — but in binary, we get to choose what a digit means.
- ▶ Let r be the number of bits per digit.
- ▶ Then:
 - ▶ Number of digits $d = \lceil b/r \rceil$
 - ▶ Each digit ranges from 0 to $2^r - 1$
 - ▶ Counting sort takes $\Theta(n + 2^r)$ per pass
 - ▶ Total time: $\Theta(d(n + 2^r)) = \Theta\left(\frac{b}{r}(n + 2^r)\right)$

Example: Sorting 32-bit Keys with Byte-wise Digits

- ▶ Let $b = 32$, $r = 8$ (one digit = 1 byte).

- ▶ Then:

- ▶ Number of digits: $d = b/r = 4$

- ▶ Digit range: 0 to $2^8 - 1 = 255$

- ▶ Each counting sort pass: $\Theta(n + 256)$

- ▶ Total cost:

$$\Theta(4(n + 256)) = \Theta(n)$$

- ▶ **Conclusion:** Efficient for fixed-width binary keys like 32-bit integers.

Why Is This Useful?

- ▶ **Tuning** r gives you flexibility:
 - ▶ Larger r : fewer passes, but larger bins
 - ▶ Smaller r : more passes, but smaller bins
- ▶ **Summary:** Lemma 8.4 gives a general formula to analyze radix sort's behavior on binary data.

Chapter 8: Sorting in Linear Time

- ▶ Chapter 8.1: Lower Bounds for Sorting
- ▶ Chapter 8.2: Counting Sort
- ▶ Chapter 8.3: Radix Sort
- ▶ **Chapter 8.4: Bucket Sort**

Bucket Sort Overview

Bucket sort assumes the input consists of n real numbers drawn **uniformly and independently** from the interval $(0, 1)$.

- ▶ Like **counting sort**, it achieves **average-case** running time of $\Theta(n)$ by making assumptions about the input.
- ▶ **Counting sort** assumes input integers lie in a small range.
- ▶ **Bucket sort** assumes inputs are randomly spread over $(0, 1)$ (uniform distribution).

Key idea: Partition the interval $(0, 1)$ into n equal-sized buckets, distribute elements into these buckets, sort within each bucket, and concatenate.

Bucket Sort: Why It Works Well

- ▶ Since inputs are **uniformly and independently** distributed:
 - ▶ Most buckets will have very few elements
 - ▶ This makes sorting inside each bucket fast
- ▶ We sort each bucket using a simple algorithm like **insertion sort**
- ▶ Finally, we **concatenate** the sorted buckets in order

Conclusion: With high probability, the total running time is $\Theta(n)$

Bucket Sort Assumptions and Structure

Assumptions:

- ▶ Input array $A[1 \dots n]$ such that each $A[i] \in [0, 1)$

Algorithm structure:

- ▶ Create an auxiliary array $B[0 \dots n - 1]$ of empty linked lists (buckets)
- ▶ For each input $A[i]$, place it in bucket:

$$B[\lfloor n \cdot A[i] \rfloor]$$

- ▶ Sort each bucket using **insertion sort**
- ▶ Scan all buckets in order and concatenate results

Example Setup: Input Array $A[1 \dots 10]$

Input values (uniform in $[0, 1)$):

$A = [0.78, 0.17, 0.39, 0.26, 0.72, 0.94, 0.21, 0.12, 0.23, 0.68]$

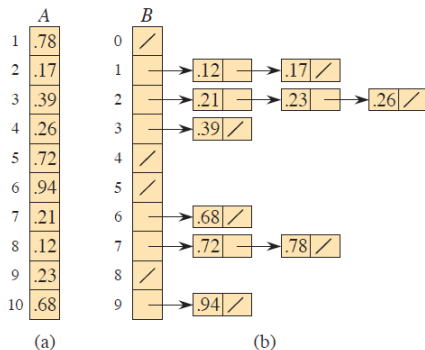
Buckets $B[0 \dots 9]$ will be created. Each value goes into:

$$\text{Bucket } \lfloor 10 \cdot A[i] \rfloor$$

For example:

- ▶ $0.78 \rightarrow B[7]$
- ▶ $0.17 \rightarrow B[1]$
- ▶ $0.39 \rightarrow B[3]$
- ▶ $0.94 \rightarrow B[9]$

Example of BUCKET-SORT ($n = 10$)



- ▶ (a) Input array $A[1 \dots 10]$
- ▶ (b) Buckets $B[0 \dots 9]$ after distributing and sorting
- ▶ Each bucket $B[i]$ holds values in the interval:

$$\left[\frac{i}{10}, \frac{i+1}{10} \right)$$

- ▶ Slashes indicate the end of each bucket
- ▶ Final output is the concatenation of all buckets in order

BUCKET-SORT Pseudocode and Why It Works

BUCKET-SORT(A, n)

```
1  let  $B[0 : n - 1]$  be a new array
2  for  $i = 0$  to  $n - 1$ 
3      make  $B[i]$  an empty list
4  for  $i = 1$  to  $n$ 
5      insert  $A[i]$  into list  $B[\lfloor n \cdot A[i] \rfloor]$ 
6  for  $i = 0$  to  $n - 1$ 
7      sort list  $B[i]$  with insertion sort
8  concatenate the lists  $B[0], B[1], \dots, B[n - 1]$  together in order
9  return the concatenated lists
```

Why it works:

- ▶ For any two elements $A[i] \leq A[j]$:
 - ▶ If they go into the same bucket \rightarrow insertion sort maintains order (line 7)
 - ▶ If they go into different buckets \rightarrow lower-indexed bucket comes first (line 8)
- ▶ So final output is fully sorted

Analyzing the Running Time of Bucket Sort

Observation: All lines **except line 7** (insertion sort) run in $O(n)$ time.

Main cost: Line 7 — insertion sort in each of the n buckets.

- ▶ Let n_i be the number of elements in bucket $B[i]$
- ▶ Insertion sort on each bucket takes $O(n_i^2)$

Total time:

$$T(n) = \Theta(n) + \sum_{i=0}^{n-1} O(n_i^2) \quad (8.1)$$

Average-Case Running Time of BUCKET-SORT

We analyze the expected time based on the input distribution (uniform over $[0, 1)$).

$$\begin{aligned}\mathbb{E}[T(n)] &= \mathbb{E} \left[\Theta(n) + \sum_{i=0}^{n-1} O(n_i^2) \right] \\ &= \Theta(n) + \sum_{i=0}^{n-1} \mathbb{E} [O(n_i^2)] \quad (\text{by linearity of expectation}) \\ &= \Theta(n) + \sum_{i=0}^{n-1} O(\mathbb{E}[n_i^2])\end{aligned}\tag{8.2}$$

We now compute $\mathbb{E}[n_i^2]$ for each bucket $i \in \{0, 1, \dots, n-1\}$

Claim: $\mathbb{E}[n_i^2] = 2 - \frac{1}{n}$ (Equation 8.3)

Each input element is equally likely to fall into any of the n buckets.

- ▶ Let n_i be the number of elements in bucket $B[i]$
- ▶ View n_i as the number of **successes** in n Bernoulli trials:
 - ▶ Success: input falls into bucket $B[i]$
 - ▶ Probability of success: $p = 1/n$, failure $q = 1 - 1/n$
 - ▶ So $n_i \sim \text{Binomial}(n, 1/n)$

Using binomial distribution properties:

$$\mathbb{E}[n_i] = np = 1 \quad \text{and} \quad \text{Var}(n_i) = npq = 1 - \frac{1}{n}$$

From identity:

$$\mathbb{E}[n_i^2] = \text{Var}(n_i) + (\mathbb{E}[n_i])^2 = \left(1 - \frac{1}{n}\right) + 1^2 = 2 - \frac{1}{n} \quad (8.3)$$

Final Result: Linear Expected Time

Now plug Equation (8.3) into Equation (8.2):

$$\begin{aligned}\mathbb{E}[T(n)] &= \Theta(n) + \sum_{i=0}^{n-1} O\left(2 - \frac{1}{n}\right) \\ &= \Theta(n) + n \cdot O\left(2 - \frac{1}{n}\right) \\ &= \Theta(n) + O(n) = \boxed{\Theta(n)}\end{aligned}$$

Conclusion: Bucket sort runs in expected linear time under uniform distribution.

Appendix - Binomial Distribution

The **binomial distribution** models the number of successes in n independent yes/no trials, each with the same probability of success.

Definition: Let $X \sim \text{Binomial}(n, p)$, where:

- ▶ n : number of trials
- ▶ p : probability of success on each trial
- ▶ $q = 1 - p$: probability of failure
- ▶ X : number of successes in n trials

Properties:

$$\mathbb{E}[X] = np, \quad \text{Var}(X) = np(1 - p)$$

Binomial distribution is used to model bucket sizes in the analysis of bucket sort.

Example: Tossing Coins

Imagine tossing a fair coin $n = 10$ times:

- ▶ Each toss has:
 - ▶ Success = heads, with probability $p = 0.5$
 - ▶ Failure = tails, with probability $q = 0.5$
- ▶ Let X be the number of heads obtained
- ▶ Then $X \sim \text{Binomial}(10, 0.5)$

Expected number of heads:

$$\mathbb{E}[X] = 10 \cdot 0.5 = 5$$

Question?