

Chapter 20. Elementary Graph Algorithms & 22. Shortest Paths

Joon Soo Yoo

June 5, 2025

Assignment


- ▶ Read §20.3, 22.0
- ▶ Problems
 - ▶ §20.3 - 5.(a), (b)

Chapter 20: Elementary Graph Algorithms

- ▶ Chapter 20.1: Representations of Graphs
- ▶ Chapter 20.2: Breadth-First Search
- ▶ **Chapter 20.3: Depth-First Search**
- ▶ Chapter 20.4: Topological Sort
- ▶ Chapter 20.5: Strongly Connected Components

What is DFS?

- ▶ DFS explores as deep as possible before backtracking.
- ▶ Recursively visits unexplored neighbors.
- ▶ Restarts from a new source if unvisited vertices remain.
- ▶ Produces a **DFS forest** instead of a single tree.



DFS Forest and Predecessor Subgraph

- ▶ Each vertex v has a predecessor $v.\pi$.
- ▶ The predecessor subgraph is:

$$G_\pi = (V, E_\pi) \quad \text{where} \quad E_\pi = \{(v.\pi, v) \mid v \in V, v.\pi \neq \text{NIL}\}$$

- ▶ G_π forms a **depth-first forest**.

Vertex Coloring in DFS

- ▶ **WHITE**: Not yet visited
- ▶ **GRAY**: Discovered, currently exploring
- ▶ **BLACK**: Fully explored
- ▶ Ensures that each vertex belongs to exactly one DFS tree

Timestamps in DFS

- ▶ Each vertex v is assigned two timestamps:
 - ▶ $v.d$ (discovery time)
 - ▶ $v.f$ (finish time)
- ▶ Timestamps help in:
 - ▶ Classifying edge types (tree, back, forward, cross)
 - ▶ Detecting cycles
 - ▶ Topological sorting

DFS Timestamps

- ▶ DFS assigns each vertex u two timestamps:
 - ▶ $u.d$: discovery time (when u is first visited)
 - ▶ $u.f$: finish time (after all of u 's neighbors are explored)
- ▶ Timestamps are integers in $[1, 2|V|]$
- ▶ Always: $u.d < u.f$

Vertex State Over Time

- ▶ Each vertex u changes state during DFS:
 - ▶ **WHITE** before $u.d$
 - ▶ **GRAY** between $u.d$ and $u.f$
 - ▶ **BLACK** after $u.f$
- ▶ These states ensure that vertices are visited and completed correctly

DFS Algorithm

DFS(G)

```
1  for each vertex  $u \in G.V$ 
2       $u.color = \text{WHITE}$ 
3       $u.\pi = \text{NIL}$ 
4   $time = 0$ 
5  for each vertex  $u \in G.V$ 
6      if  $u.color == \text{WHITE}$ 
7          DFS-VISIT( $G, u$ )
```

DFS Visit Algorithm

DFS-VISIT(G, u)

```
1   $time = time + 1$                 // white vertex  $u$  has just been discovered
2   $u.d = time$ 
3   $u.color = \text{GRAY}$ 
4  for each vertex  $v$  in  $G.Adj[u]$  // explore each edge  $(u, v)$ 
5      if  $v.color == \text{WHITE}$ 
6           $v.\pi = u$ 
7          DFS-VISIT( $G, v$ )
8   $time = time + 1$ 
9   $u.f = time$ 
10  $u.color = \text{BLACK}$              // blacken  $u$ ; it is finished
```

Why DFS Order Matters

- ▶ DFS depends on:
 - ▶ The order of vertices in $G.V$
 - ▶ The order of neighbors in each adjacency list $\text{Adj}[u]$
- ▶ This affects:
 - ▶ The DFS forest structure
 - ▶ Timestamps (d, f)
 - ▶ Edge classifications (tree, back, forward, cross)
- ▶ But all are valid DFS results

DFS-VISIT Calls and Edge Scanning

- ▶ DFS-VISIT(G, u) is called exactly once per vertex:

$$\text{Total calls to DFS-VISIT} = |V|$$

- ▶ Each call explores neighbors:

$$\text{for each } v \in \text{Adj}[u]$$

- ▶ Total adjacency scans:

$$\sum_{u \in V} |\text{Adj}[u]|$$

- ▶ $= |E|$ for directed graphs
- ▶ $= 2|E|$ for undirected graphs

Final Running Time of DFS


- ▶ Initialization: $\Theta(|V|)$
- ▶ DFS-VISIT calls: $\Theta(|V|)$
- ▶ Neighbor scans: $\Theta(|E|)$

$$\Theta(|V| + |E|)$$


Classifying Edges in DFS

- ▶ DFS reveals graph structure by classifying edges:
- ▶ **Tree edge:** (u, v) discovers v for the first time
- ▶ **Back edge:** (u, v) goes to an ancestor of u (or a self-loop)
- ▶ **Forward edge:** (u, v) goes to a proper descendant of u
- ▶ **Cross edge:** all other edges (between DFS trees or unrelated nodes)


DFS Diagram




(a)




(b)




(c)




(d)




(e)




(f)




(g)




(h)




(i)




(j)




(k)




(l)




(m)



(n)



(o)



(p)

Classification by Color During DFS

When DFS explores edge (u, v) :

- ▶ v is **WHITE** \Rightarrow Tree edge
- ▶ v is **GRAY** \Rightarrow Back edge (ancestor)
- ▶ v is **BLACK**:
 - ▶ If $u.d < v.d$ and $v.f < u.f \Rightarrow$ Forward edge
 - ▶ If $v.f < u.d \Rightarrow$ Cross edge

Why Edge Classification Matters

- ▶ Back edges indicate **cycles**
- ▶ Forward and cross edges only occur in **directed** graphs
- ▶ DFS has enough info to classify edges:
 - ▶ via vertex colors at traversal time
 - ▶ via discovery/finish times (d, f)
- ▶ Helps in graph analysis

Chapter 22: Single-Source Shortest Paths

- ▶ Chapter 22.1: The Bellman-Ford Algorithm
- ▶ Chapter 22.2: Single-Source Shortest Paths in Directed Acyclic Graphs
- ▶ Chapter 22.3: Dijkstra's Algorithm
- ▶ Chapter 22.4: Difference Constraints and Shortest Paths
- ▶ Chapter 22.5: Proofs of Shortest-Paths Properties

Finding Shortest Routes – Real-World Motivation

- ▶ Drive from **NY** to **CA**
- ▶ GPS models:
 - ▶ Intersections \rightarrow vertices
 - ▶ Roads \rightarrow directed edges
 - ▶ Distances \rightarrow edge weights
- ▶ Enumerating all routes is infeasible

Shortest-Path Problem Setup

Input:

- ▶ Directed graph $G = (V, E)$
- ▶ Weight function $w : E \rightarrow \mathbb{R}$

Path weight:

$$w(p) = \sum_{i=1}^k w(v_{i-1}, v_i)$$

Shortest-path weight:

$$\delta(u, v) = \begin{cases} \min\{w(p) : u \rightsquigarrow v\} & \text{if path exists} \\ \infty & \text{otherwise} \end{cases}$$

Edge Weights – More Than Just Distance

- ▶ Weights can represent:
 - ▶ Travel time
 - ▶ Monetary cost
 - ▶ Penalty or loss
 - ▶ Any additive metric
- ▶ Goal: Minimize total cost along the path

Optimal Substructure of Shortest Paths

- ▶ Many shortest-path algorithms rely on the idea of **optimal substructure**.
- ▶ That is, a shortest path between two vertices contains other shortest paths within it.
- ▶ This property enables the use of:
 - ▶ **Greedy algorithms**, e.g., **Dijkstra's algorithm** (Section 22.3)
 - ▶ **Dynamic programming**, e.g., **Floyd-Warshall algorithm** (Section 23.2)
- ▶ Similar principle is also used in the Edmonds-Karp algorithm (Chapter 24, max flow).

Lemma 22.1 formalizes this optimal-substructure property.

Lemma 22.1: Optimal Substructure of Shortest Paths

Lemma

Given a weighted, directed graph $G = (V, E)$ with weight function $w : E \rightarrow \mathbb{R}$, let

$$p = \langle v_0, v_1, \dots, v_k \rangle$$

be a shortest path from v_0 to v_k . Then for any $0 \leq i \leq j \leq k$, the subpath

$$p_{ij} = \langle v_i, v_{i+1}, \dots, v_j \rangle$$

is also a shortest path from v_i to v_j .

Proof of Lemma 22.1

Proof Sketch:

- ▶ Suppose p_{ij} is not a shortest path.
- ▶ Then there exists another path p'_{ij} with lower weight:
 $w(p'_{ij}) < w(p_{ij})$.
- ▶ Replacing p_{ij} with p'_{ij} in p gives a path from v_0 to v_k with smaller total weight.
- ▶ This contradicts the assumption that p is a shortest path.

Negative-Weight Edges in Shortest Paths



- ▶ Some graphs contain edges with negative weights.
- ▶ These do not always cause problems:
 - ▶ If there are **no negative-weight cycles reachable from the source s** ,

$\delta(s, v)$ is well-defined for all $v \in V$

even if $\delta(s, v) < 0$.

- ▶ But if a **negative-weight cycle is reachable from s** :
 - ▶ You can cycle repeatedly and lower path weight indefinitely.
 - ▶ So: $\delta(s, v) = -\infty$ for any v reachable via that cycle.

Negative-Weight Edges Example



Example: Negative Cycles in Action (Fig. 22.1)

- ▶ Example paths:
 - ▶ $s \rightarrow a$: weight = 3, so $\delta(s, a) = 3$
 - ▶ $s \rightarrow a \rightarrow b$: $3 + (-4) = -1$, so $\delta(s, b) = -1$
 - ▶ $s \rightarrow c$: directly with weight = 5
 - ▶ $c \rightarrow d \rightarrow c$: forms a **positive-weight cycle**, so no problem.
- ▶ Cycle $e \rightarrow f \rightarrow e$: weight = $3 + (-6) = -3 \rightarrow$
negative-weight cycle reachable from s
- ▶ Now we can:
 - ▶ Loop arbitrarily through $e \leftrightarrow f$ to decrease cost
 - ▶ Then exit to g : $s \rightarrow e \rightarrow f \rightarrow g$
- ▶ So: $\delta(s, e) = \delta(s, f) = \delta(s, g) = -\infty$

Impact on Algorithms

- ▶ **Dijkstra's algorithm:**

- ▶ Assumes all edge weights ≥ 0
- ▶ Will fail or produce incorrect results with negative-weight edges

- ▶ **Bellman-Ford algorithm:**

- ▶ Handles negative weights
- ▶ Produces correct results as long as no negative-weight cycle is reachable from s
- ▶ Can **detect negative-weight cycles**

Summary: Negative-weight edges are okay, but negative-weight cycles are dangerous.

Cycles in Shortest Paths

- ▶ Can a shortest path contain a cycle?
- ▶ **Negative-weight cycle:** Repeating it reduces the total weight \rightarrow no well-defined shortest path.
- ▶ **Positive-weight cycle:** Removing the cycle makes the path shorter:

$$w(p') = w(p) - w(c) < w(p)$$

So p wasn't the shortest path.

- ▶ **Conclusion:** We can always assume that a shortest path is a **simple path** (i.e., no cycles).

Representing Shortest Paths

Why store paths, not just distances?

- ▶ Knowing only the shortest-path *distance* is often not enough.
- ▶ Applications (e.g., GPS) require the actual *path*.

Storing paths with predecessors:

- ▶ For each vertex v , store $v.\pi$ (predecessor of v on shortest path).
- ▶ Use PRINT-PATH(G, s, v) to reconstruct the path from s to v .

Predecessor subgraph:

$$V_\pi = \{v \in V \mid v.\pi \neq \text{NIL}\} \cup \{s\}, \quad E_\pi = \{(v.\pi, v) \in E \mid v \in V_\pi \setminus \{s\}\}$$

Shortest-Paths Tree

Definition: A *shortest-paths tree* from source s is a directed subgraph $G' = (V', E')$ such that:

1. V' is the set of vertices reachable from s
2. G' is a rooted tree with root s
3. For all $v \in V'$, the unique path from s to v in G' is a shortest path in G

Key Properties:

- ▶ Edge weights are used (unlike BFS trees which use hop-count).
- ▶ Shortest paths (and trees) may not be unique.

Shortest-Paths Tree Example




Figure 22.2 (a) A weighted, directed graph with shortest-path weights from source s . (b) The blue edges form a shortest-paths tree rooted at the source s . (c) Another shortest-paths tree with the same root.

Relaxation in Shortest-Paths Algorithms

Core Concept: Relaxation

- ▶ Each vertex v maintains a **shortest-path estimate** $v.d$
 - ▶ $v.d$ is an upper bound on the weight of the shortest path from source s to v
- ▶ The goal is to iteratively **reduce** $v.d$ to the correct shortest-path weight $\delta(s, v)$

Initialization:

INITIALIZE-SINGLE-SOURCE(G, s)

```
1  for each vertex  $v \in G.V$ 
2       $v.d = \infty$ 
3       $v.\pi = \text{NIL}$ 
4   $s.d = 0$ 
```

The RELAX Operation

Relaxing an edge (u, v) with weight $w(u, v)$:

```
RELAX( $u, v, w$ )
```

```
1  if  $v.d > u.d + w(u, v)$   
2       $v.d = u.d + w(u, v)$   
3       $v.\pi = u$ 
```

Effect:

- ▶ Updates $v.d$ if a better path through u is found
- ▶ Updates $v.\pi$ to point to u

The RELAX Example




Figure 22.3 Relaxing an edge (u, v) with weight $w(u, v) = 2$. The shortest-path estimate of each vertex appears within the vertex. **(a)** Because $v.d > u.d + w(u, v)$ prior to relaxation, the value of $v.d$ decreases. **(b)** Since we have $v.d \leq u.d + w(u, v)$ before relaxing the edge, the relaxation step leaves $v.d$ unchanged.

Question?