

# Chapter 20.3 Dijkstra's Algorithm & Chapter 34 NP-Completeness

Joon Soo Yoo

June 10, 2025

# Assignment

- ▶ Read §22.3, §34.1
- ▶ Problems
  - ▶ §22.3 - 2

## Chapter 22: Single-Source Shortest Paths

- ▶ Chapter 22.1: The Bellman-Ford Algorithm
- ▶ Chapter 22.2: Single-Source Shortest Paths in Directed Acyclic Graphs
- ▶ **Chapter 22.3: Dijkstra's Algorithm**
- ▶ Chapter 22.4: Difference Constraints and Shortest Paths
- ▶ Chapter 22.5: Proofs of Shortest-Paths Properties

# Relaxation in Shortest-Paths Algorithms

## Core Concept: Relaxation

- ▶ Each vertex  $v$  maintains a **shortest-path estimate**  $v.d$ 
  - ▶  $v.d$  is an upper bound on the weight of the shortest path from source  $s$  to  $v$
- ▶ The goal is to iteratively **reduce**  $v.d$  to the correct shortest-path weight  $\delta(s, v)$

## Initialization:

INITIALIZE-SINGLE-SOURCE( $G, s$ )

```
1  for each vertex  $v \in G.V$ 
2       $v.d = \infty$ 
3       $v.\pi = \text{NIL}$ 
4   $s.d = 0$ 
```

# The RELAX Operation

**Relaxing an edge  $(u, v)$  with weight  $w(u, v)$ :**


```
RELAX( $u, v, w$ )
```

```
1  if  $v.d > u.d + w(u, v)$   
2       $v.d = u.d + w(u, v)$   
3       $v.\pi = u$ 
```

**Effect:**

- ▶ Updates  $v.d$  if a better path through  $u$  is found
- ▶ Updates  $v.\pi$  to point to  $u$

# The RELAX Example



**Figure 22.3** Relaxing an edge  $(u, v)$  with weight  $w(u, v) = 2$ . The shortest-path estimate of each vertex appears within the vertex. **(a)** Because  $v.d > u.d + w(u, v)$  prior to relaxation, the value of  $v.d$  decreases. **(b)** Since we have  $v.d \leq u.d + w(u, v)$  before relaxing the edge, the relaxation step leaves  $v.d$  unchanged.

# Dijkstra's Algorithm: Overview

**Goal:** Find the shortest paths from a source vertex  $s$  to all other vertices in a weighted, directed graph  $G = (V, E)$ , with  $w(u, v) \geq 0$ .

## Key Concepts:

- ▶ Generalizes BFS to weighted graphs
- ▶ Uses a **greedy strategy** to iteratively build a shortest-paths tree
- ▶ Maintains:
  - ▶ Set  $S$ : vertices whose shortest-path weights are known
  - ▶ Min-priority queue  $Q = V - S$


# Dijkstra's Algorithm: Pseudocode

DIJKSTRA( $G, w, s$ )


```
1  INITIALIZE-SINGLE-SOURCE( $G, s$ )
2   $S = \emptyset$ 
3   $Q = \emptyset$ 
4  for each vertex  $u \in G.V$ 
5      INSERT( $Q, u$ )
6  while  $Q \neq \emptyset$ 
7       $u = \text{EXTRACT-MIN}(Q)$ 
8       $S = S \cup \{u\}$ 
9      for each vertex  $v$  in  $G.Adj[u]$ 
10         RELAX( $u, v, w$ )
11         if the call of RELAX decreased  $v.d$ 
12             DECREASE-KEY( $Q, v, v.d$ )
```




# Dijkstra's Algorithm: Diagram




(a)




(b)




(c)



(d)



(e)



(f)

# Dijkstra's Algorithm: Step-by-Step (I)

**Input:** Directed graph  $G = (V, E)$  with  $w(u, v) \geq 0$  for all  $(u, v) \in E$  and source  $s$

## Initialization:

1. For each  $v \in V$ : set  $v.d \leftarrow \infty$ ,  $v.\pi \leftarrow \text{NIL}$ ; then set  $s.d \leftarrow 0$
2. Set  $S \leftarrow \emptyset$  (vertices with finalized shortest-path weights)
3. Insert all  $v \in V$  into a min-priority queue  $Q$ , keyed by  $v.d$

## Main Loop:

- ▶ While  $Q \neq \emptyset$ :
  - ▶  $u \leftarrow \text{EXTRACT-MIN}(Q)$
  - ▶ Add  $u$  to  $S$
  - ▶ For each  $v \in \text{Adj}[u]$ , perform  $\text{RELAX}(u, v, w)$
  - ▶ If  $v.d$  changes, call  $\text{DECREASE-KEY}(Q, v, v.d)$

# Dijkstra's Algorithm: Key Properties (II)

## Greedy Strategy:

- ▶ Always selects vertex  $u \in V - S$  with the smallest  $u.d$
- ▶ Once  $u$  is added to  $S$ ,  $u.d = \delta(s, u)$  is guaranteed

## Invariant:

- ▶  $Q = V - S$  at the start of each iteration
- ▶ Each vertex is extracted from  $Q$  exactly once
- ▶ Each edge  $(u, v)$  is relaxed at most once

**Output:** For all  $v \in V$ ,  $v.d = \delta(s, v)$  and  $v.\pi$  gives the predecessor on the shortest path

# Correctness of Dijkstra's Algorithm

## Theorem (Theorem 22.6)

*Let  $G = (V, E)$  be a weighted, directed graph with nonnegative edge weights  $w(u, v) \geq 0$  and source vertex  $s \in V$ .*

*Then, after Dijkstra's algorithm terminates, for every vertex  $u \in V$ , we have:*

$$u.d = \delta(s, u)$$

*That is, the algorithm correctly computes the shortest-path distance from  $s$  to  $u$ .*

## Corollary (Corollary 22.7)

*The predecessor subgraph  $G_\pi$  forms a shortest-paths tree rooted at  $s$ .*

## Proof Sketch of Theorem 22.6 (1/2)

**Strategy:** Use induction on the number of iterations of the while loop (i.e., size of  $S$ ).

**Base case:**

- ▶ When  $S = \emptyset$ , no vertex is added, so the claim trivially holds.
- ▶ When  $S = \{s\}$ , we have  $s.d = 0 = \delta(s, s)$ .

**Inductive step:**

- ▶ Assume for all  $v \in S$ ,  $v.d = \delta(s, v)$ .
- ▶ Let  $u$  be the next vertex extracted from  $Q = V - S$  (minimum  $d$ ).
- ▶ Show:  $u.d = \delta(s, u)$ .

**Key idea:**

- ▶ Let  $y$  be the first vertex on a shortest path from  $s$  to  $u$  that is not in  $S$ .
- ▶ Let  $x \in S$  be the predecessor of  $y$  on this path.

## Proof Sketch of Theorem 22.6 (2/2)

### Argument:

- ▶ From the inductive hypothesis:  $x.d = \delta(s, x)$ .
- ▶ Relaxation of edge  $(x, y)$  ensures:  $y.d = \delta(s, y)$  at the time  $x$  was added to  $S$ .
- ▶ Because  $u$  was chosen as  $\text{EXTRACT-MIN}(Q)$  and  $y$  is on the path to  $u$ , we have:

$$\delta(s, y) \leq \delta(s, u) \leq u.d \leq y.d$$

and since  $y.d = \delta(s, y)$ , all inequalities collapse:

$$u.d = \delta(s, u)$$

**Conclusion:** By induction, when the algorithm terminates,  $u.d = \delta(s, u)$  for all  $u \in V$ .

# Time Complexity of Dijkstra's Algorithm (Simple Array Implementation)

**Assumption:** Min-priority queue implemented using an *unsorted array*.

## Operation Costs:

- ▶  $\text{INSERT}(v): \mathcal{O}(1)$
- ▶  $\text{EXTRACT-MIN}(): \mathcal{O}(|V|)$
- ▶  $\text{DECREASE-KEY}(v): \mathcal{O}(1)$

## Total Number of Calls:

- ▶  $|V|$  calls to  $\text{INSERT}$  and  $\text{EXTRACT-MIN}$
- ▶ At most  $|E|$  calls to  $\text{DECREASE-KEY}$  ( $\because$  every edge relaxes once)

## Total Time:

$$\mathcal{O}(|V| \cdot |V| + |E| \cdot 1) = \mathcal{O}(|V|^2 + |E|) = \mathcal{O}(|V|^2)$$

**Use case:** Acceptable for *dense graphs*

# Time Complexity of Dijkstra's Algorithm (Binary Min-Heap Implementation)

**Assumption:** Min-priority queue implemented using a *binary min-heap* with vertex-to-index mapping.

## Operation Costs:

- ▶  $\text{INSERT}(v)$ :  $\mathcal{O}(\log |V|)$  — insert into heap
- ▶  $\text{EXTRACT-MIN}()$ :  $\mathcal{O}(\log |V|)$  — remove min and heapify
- ▶  $\text{DECREASE-KEY}(v)$ :  $\mathcal{O}(\log |V|)$  — heapify

## Total Number of Calls:

- ▶  $|V|$  calls to  $\text{INSERT}$  and  $\text{EXTRACT-MIN}$
- ▶ At most  $|E|$  calls to  $\text{DECREASE-KEY}$

## Total Time:

$$\mathcal{O}(|V| \log |V| + |E| \log |V|) = \mathcal{O}((|V| + |E|) \log |V|)$$

**Use case:** Efficient for *sparse graphs*



# Chapter 34: NP-Completeness


- ▶ Chapter 34.1: Polynomial Time
- ▶ Chapter 34.2: Polynomial-Time Verification
- ▶ Chapter 34.3: NP-Completeness and Reducibility
- ▶ Chapter 34.4: NP-Completeness Proofs
- ▶ Chapter 34.5: NP-Complete Problems

# P, NP, and NP-Completeness: Informal Overview

- ▶ **P:** Problems solvable in polynomial time (e.g.,  $O(n^k)$  for some constant  $k$ ).
- ▶ **NP:** Problems verifiable in polynomial time. Given a solution (*certificate*), it can be verified quickly.
- ▶ **NP-Complete (NPC):** Problems in NP that are as hard as any other problem in NP.

If any NP-complete problem is in P, then **P = NP**.

# Examples of NP Verifiability



## Hamiltonian Cycle:

- ▶ Input: Graph  $G = (V, E)$
- ▶ Certificate: Sequence  $\langle v_1, v_2, \dots, v_{|V|} \rangle$
- ▶ Verify: Each node appears once; all edges exist;  $v_{|V|} \rightarrow v_1$

# P vs NP

- ▶ Every problem in **P** is in **NP**.
- ▶ Open question: Is **P**  $\subsetneq$  **NP**?
- ▶ If any NP-complete problem is in **P**, then **P** = **NP**.

Most researchers believe **P**  $\neq$  **NP**.

# What is NP-Completeness?

A problem  $Q$  is NP-complete if:

1.  $Q \in \mathbf{NP}$  (verifiable in poly time)
2. Every problem in  $\mathbf{NP}$  can be *reduced* to  $Q$  in poly time

If any NP-complete problem is solvable in poly time, all of NP is.

# Why NP-Completeness Matters

- ▶ Proving NP-completeness shows likely intractability
- ▶ Better to design:
  - ▶ Approximation algorithms
  - ▶ Heuristics
  - ▶ Algorithms for special cases
- ▶ Many natural problems turn out to be NP-complete

# Optimization vs Decision Problems

**Optimization:** Find best solution (e.g., shortest path, min cost).

**Decision:** Yes/No form. E.g.,

- ▶ **PATH:** Is there a path from  $u$  to  $v$  with  $\leq k$  edges?


We study NP-completeness through decision problems.

# Reductions: Relating Problem Hardness

- ▶ Given problems  $A$  and  $B$
- ▶ If you can reduce  $A$  to  $B$  in poly time:
  - ▶  $A \leq_p B$ :  $A$  is *no harder* than  $B$
- ▶ Use known hard problems to show new problems are hard



# Reduction Example: Visual



**Goal:** Show that problem  $A$  is solvable in polynomial time.

**Algorithm  $A'$  for solving  $A$ :**

1. Input: instance  $\alpha$  of problem  $A$
2. Compute  $\beta = f(\alpha)$  in polynomial time
3. Run the known poly-time algorithm for  $B$  on  $\beta$
4. Output the same answer for  $\alpha$

**Conclusion:**  $A$  is solvable in polynomial time.

# Proving Intractability via Reductions

Suppose:

- ▶ Problem A is known to be hard (no poly algorithm)
- ▶ You reduce A to B (i.e.,  $A \leq_p B$ )

Then:

- ▶ If B had a poly-time algorithm, so would A — contradiction!
- ▶ So B must also be hard

This is how we prove NP-completeness.

# Chapter 34: NP-Completeness

- ▶ **Chapter 34.1: Polynomial Time**
- ▶ Chapter 34.2: Polynomial-Time Verification
- ▶ Chapter 34.3: NP-Completeness and Reducibility
- ▶ Chapter 34.4: NP-Completeness Proofs
- ▶ Chapter 34.5: NP-Complete Problems

# Abstract Problems: Formal Definition

**Abstract problem**  $Q$  is defined as a **binary relation** between:

- ▶ A set  $I$  of **problem instances**, and
- ▶ A set  $S$  of **problem solutions**

**Example:** SHORTEST-PATH

- ▶ Instance:  $(G, u, v)$  — a graph and two vertices
- ▶ Solution: A path from  $u$  to  $v$  (or empty if none exists)
- ▶ Relation: Matches each instance with a valid shortest path

*Note:* There may be multiple valid solutions — shortest paths are not necessarily unique.

# From Abstract Problems to Decision Problems

## Why focus on decision problems?

- ▶ NP-completeness theory deals with problems having **yes/no answers**
- ▶ These are easier to define and analyze formally

**A decision problem** is a function:

$$f : I \rightarrow \{0, 1\}$$

- ▶  $f(i) = 1$  (yes) if the instance satisfies the condition
- ▶  $f(i) = 0$  (no) otherwise

## Example: PATH

- ▶ Input:  $(G, u, v, k)$
- ▶ Output: 1 if there exists a path from  $u$  to  $v$  with at most  $k$  edges; otherwise 0

# Optimization vs. Decision Problems

## Many problems are optimization problems:

- ▶ Find the shortest path, largest flow, minimum cut, etc.

## But we can often transform them into decision problems:

- ▶ Ask: “*Is there a solution of cost  $\leq k$ ?*”
- ▶ This transformation is key in NP-completeness proofs

## Why this works:

- ▶ The decision version is *no harder* than the optimization version
- ▶ We can recover optimal values using repeated queries (e.g., binary search)

# Optimization vs. Decision: Example

## Problem: SHORTEST PATH

### Optimization Version

- ▶ Given a graph  $G$  and vertices  $u, v$
- ▶ **Goal:** Find the *length of the shortest path* from  $u$  to  $v$

### Decision Version

- ▶ Given a graph  $G$ , vertices  $u, v$ , and integer  $k$
- ▶ **Question:** Is there a path from  $u$  to  $v$  with length  $\leq k$ ?
- ▶ Answer is **Yes/No**

*Both versions deal with the same underlying problem, but the decision version is used for NP-completeness analysis.*

# What is an Encoding?

Encoding is how we represent abstract problem instances as binary strings that a computer can understand.

**Formal Definition:** An *encoding* is a function:

$$e : S \rightarrow \{0, 1\}^*$$

that maps abstract objects  $S$  to binary strings.

**Examples:**

- ▶  $e(17) = 10001$
- ▶  $e(A) = 01000001$  (ASCII)
- ▶ A graph  $G$ : encoded as adjacency list, matrix, or edge list

Compound objects (sets, graphs, programs) are encoded by combining the encodings of their parts.



# From Abstract to Concrete Problems

To run on a computer, an abstract decision problem must be encoded.

## Concrete Problem:

- ▶ Input: a binary string  $x \in \{0, 1\}^*$
- ▶ Output:  $f(x) \in \{0, 1\}$

**Size of instance:**  $n = |x|$  (length of binary string)

**Polynomial-time algorithm:** solves any input of length  $n$  in  $O(n^k)$  for some constant  $k$ .

# Encoding Can Affect Complexity

Encoding affects whether an algorithm appears to run in polynomial or exponential time.

## **Example: Integer $k$ given as input**

- ▶ Unary encoding:  $111 \dots 1$  (length  $n = k$ )  $\rightarrow$  runtime  $\Theta(k) = \Theta(n) \rightarrow$  polynomial
- ▶ Binary encoding:  $10001$  (length  $n = \lfloor \log_2 k \rfloor + 1$ )  $\rightarrow$  runtime  $\Theta(k) = \Theta(2^n) \rightarrow$  exponential

**Conclusion:** The *choice of encoding* can drastically change the apparent complexity.

# Polynomially Related Encodings

**Goal:** We want our notion of “efficient” (polynomial-time) to be encoding-independent.

**Definition:** Encodings  $e_1$  and  $e_2$  are *polynomially related* if:

- ▶ There exist functions  $f_{12}, f_{21}$  computable in polynomial time
- ▶ Such that  $f_{12}(e_1(i)) = e_2(i)$  and  $f_{21}(e_2(i)) = e_1(i)$  for all instances  $i$

This ensures we can convert between encodings without changing the complexity class.

## Lemma 34.1: Encoding-Invariant Polynomial Time

**Lemma:** Let  $Q$  be an abstract decision problem.  
If encodings  $e_1$  and  $e_2$  are polynomially related, then:

$$e_1(Q) \in \mathbf{P} \iff e_2(Q) \in \mathbf{P}$$

### Implication:

- ▶ Polynomial-time solvability doesn't depend on which *reasonable* encoding you choose
- ▶ We can now safely talk about abstract problems being in  $\mathbf{P}$ , assuming a standard encoding

# Practical Takeaways on Encodings

- ▶ You must encode instances as binary strings to analyze complexity formally
- ▶ Bad encodings (e.g., unary) can make easy problems look hard
- ▶ As long as we use **reasonable encodings** (e.g., binary, base-3, ASCII), complexity classes are preserved
- ▶ We usually assume a *standard encoding* (e.g.,  $\langle G \rangle$  for a graph)

This allows us to focus on the *problem itself*, not the details of its representation.

# A Formal-Language Framework for Decision Problems

**Alphabet:** A finite set of symbols (denoted  $\Sigma$ )

**Language:** A set of strings over  $\Sigma$  (i.e.,  $L \subseteq \Sigma^*$ )

**Examples:**

- ▶  $\Sigma = \{0, 1\}$
- ▶  $L = \{10, 11, 101, 111, 10001, \dots\}$  (binary encodings of prime numbers)

**Useful notation:**

- ▶  $\epsilon$ : empty string
- ▶  $\Sigma^*$ : set of all binary strings

# Viewing Decision Problems as Languages

**Decision problem**  $Q : \Sigma^* \rightarrow \{0, 1\}$  maps binary strings to yes/no.

We define the associated language:

$$L_Q = \{x \in \Sigma^* \mid Q(x) = 1\}$$

This is the set of inputs for which the answer is **yes**.

**Example: PATH**

- ▶ Input:  $x = \langle G, u, v, k \rangle$  (encoded as a binary string)
- ▶  $x \in L_{\text{PATH}}$  if  $G$  contains a path from  $u$  to  $v$  of length  $\leq k$

This lets us study decision problems as formal languages, enabling complexity class definitions.

# Accepting vs. Deciding a Language

**Algorithm  $A$  accepts  $x$**  if  $A(x) = 1$  (i.e., it outputs yes).

**Language accepted by  $A$ :**

$$L = \{x \in \Sigma^* \mid A(x) = 1\}$$

**But:**  $A$  might not halt on  $x \notin L$  (it might loop forever).

**Algorithm  $A$  decides  $L$**  if:

- ▶  $A(x) = 1$  for  $x \in L$
- ▶  $A(x) = 0$  for  $x \notin L$



# Polynomial-Time Acceptance vs. Decision

## Accepted in polynomial time:

- ▶  $A$  outputs 1 for  $x \in L$  in  $O(n^k)$  time
- ▶ No guarantee for  $x \notin L$  (may loop forever)

## Decided in polynomial time:

- ▶  $A$  halts on all inputs
- ▶  $A(x) = 1$  if  $x \in L$ , and  $A(x) = 0$  otherwise

**Example:** PATH can be both accepted and decided in polynomial time

# Defining **P** via Formal Languages

## Complexity Class **P**:

$\mathbf{P} = \{L \subseteq \Sigma^* \mid \text{there exists an algorithm that decides } L \text{ in } O(n^k) \text{ time}\}$

## Theorem 34.2:

- ▶ **P** is the set of languages that are *accepted* by a polynomial-time algorithm
- ▶ Proof: Any polynomial-time accepting algorithm can be converted to a deciding algorithm via bounded simulation

# Question?